

# Deep generative modeling: Latent Variable Models

Jakub M. Tomczak  
Deep Learning 2020

# INTRODUCTION

# IS GENERATIVE MODELING IMPORTANT?

We learn a neural network to classify images:

# IS GENERATIVE MODELING IMPORTANT?

We learn a neural network to classify images:



# IS GENERATIVE MODELING IMPORTANT?

We learn a neural network to classify images:

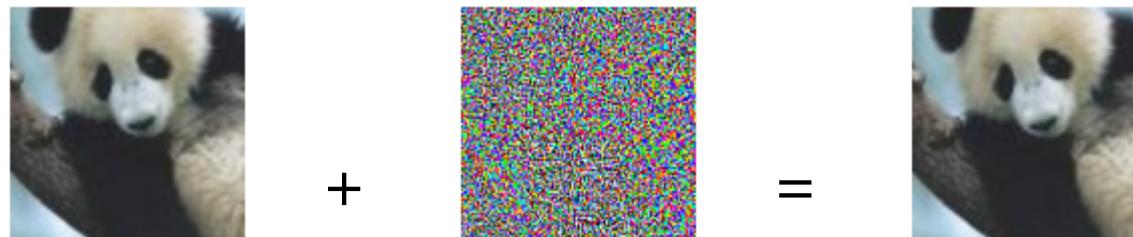


$$p(\text{panda} | x) = 0.99$$

...

# IS GENERATIVE MODELING IMPORTANT?

We learn a neural network to classify images:



$$p(\text{panda} | x) = 0.99$$

...

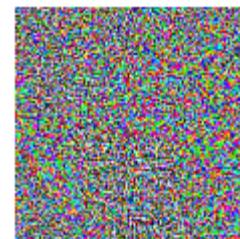
noise

# IS GENERATIVE MODELING IMPORTANT?

We learn a neural network to classify images:



+



=



$p(\text{panda}|x)=0.99$

...

noise

$p(\text{panda}|x)=0.01$

...

$p(\text{dog}|x)=0.9$

# IS GENERATIVE MODELING IMPORTANT?

We learn a neural network to classify images:



$$p(\text{panda}|x)=0.99$$

...

noise

$$p(\text{panda}|x)=0.01$$

...

$$p(\text{dog}|x)=0.9$$

There is no semantic understanding of images.

# IS GENERATIVE MODELING IMPORTANT?

This simple example shows that:

- A discriminative model is (probably) **not enough**.
- We need a notion of **uncertainty**.
- We need to **understand** the reality.

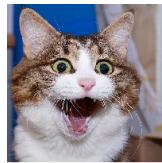
# IS GENERATIVE MODELING IMPORTANT?

This simple example shows that:

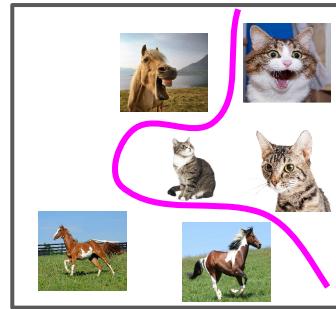
- A discriminative model is (probably) **not enough**.
- We need a notion of **uncertainty**.
- We need to **understand** the reality.

A possible solution is **generative modeling**.

# IS GENERATIVE MODELING IMPORTANT?

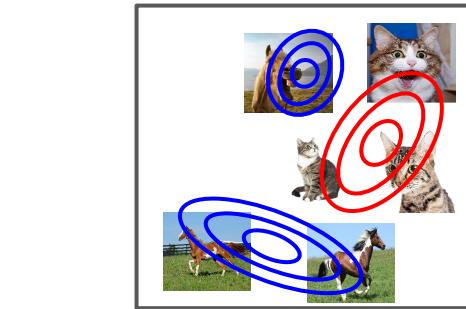
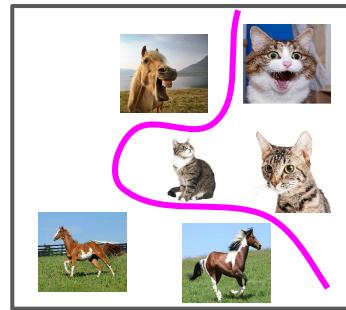
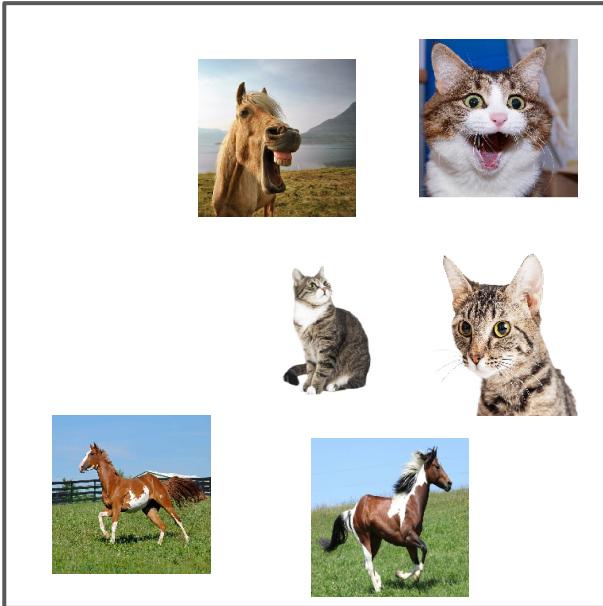


# IS GENERATIVE MODELING IMPORTANT?

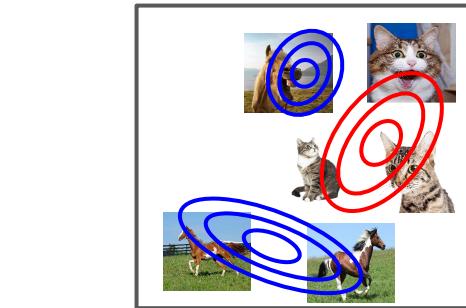
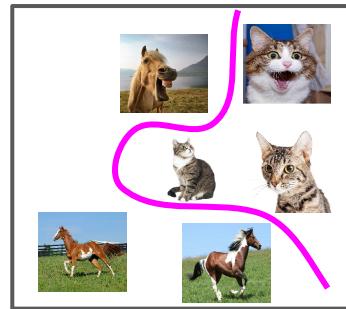


$$p_{\theta}(y|x)$$

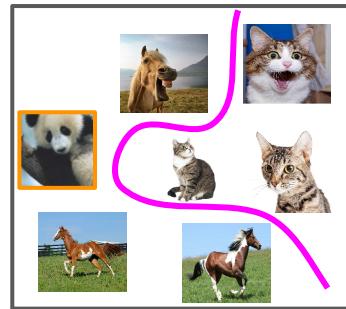
# IS GENERATIVE MODELING IMPORTANT?



# IS GENERATIVE MODELING IMPORTANT?

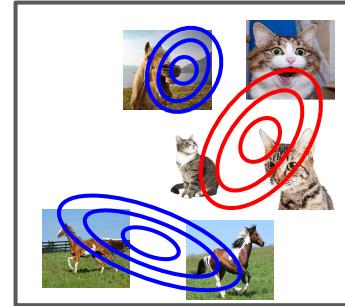


# IS GENERATIVE MODELING IMPORTANT?



$$p_{\theta}(y|x)$$

**High probability  
of a horse.  
=**  
**Highly probable  
decision!**



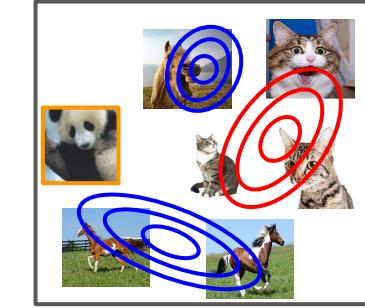
$$p_{\theta}(x, y) = p_{\theta}(y|x) p_{\theta}(x)$$

# IS GENERATIVE MODELING IMPORTANT?



$$p_{\theta}(y|x)$$

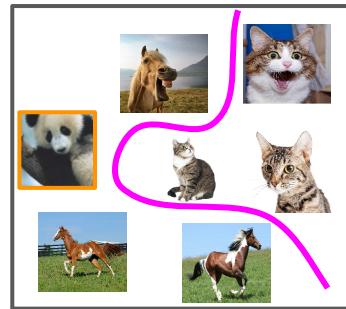
**High probability  
of a horse.  
=**  
**Highly probable  
decision!**



$$p_{\theta}(x,y) = p_{\theta}(y|x) \cdot p_{\theta}(x)$$

**High probability of  
a horse.  
x**  
**Low probability of  
the object  
=**  
**Uncertain  
decision!**

# IS GENERATIVE MODELING IMPORTANT?

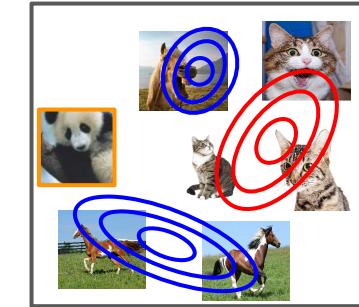


$$p_\theta(y|x)$$

**High probability  
of a horse.**

=

**Highly probable  
decision!**



$$p_\theta(x, y) = p_\theta(y|x) \circledcirc p_\theta(x)$$

**High probability of  
a horse.**

$\times$

**Low probability of  
the object**

=

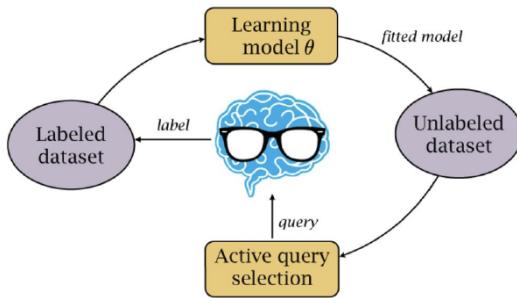
**Uncertain  
decision!**

# WHERE DO WE USE DEEP GENERATIVE MODELING?

“ i want to talk to you . ”  
“ i want to be with you . ”  
“ i do n’t want to be with you . ”  
“ i do n’t want to be with you . ”  
she did n’t want to be with him .

he was silent for a long moment .  
he was silent for a moment .  
it was quiet for a moment .  
it was dark and cold .  
there was a pause .  
it was my turn .

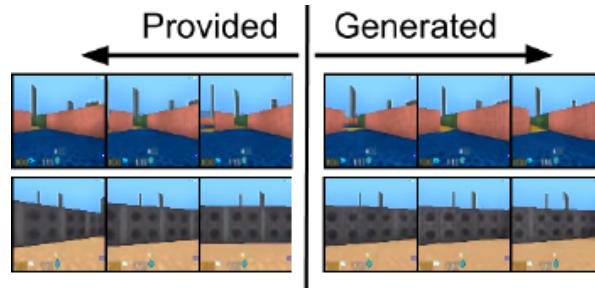
## Text analysis



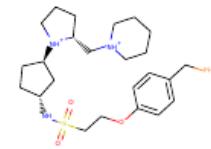
## Active Learning



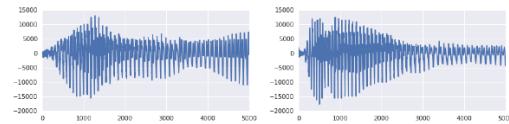
## Image analysis



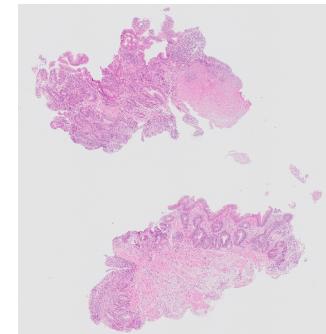
## Reinforcement Learning



## Graph analysis

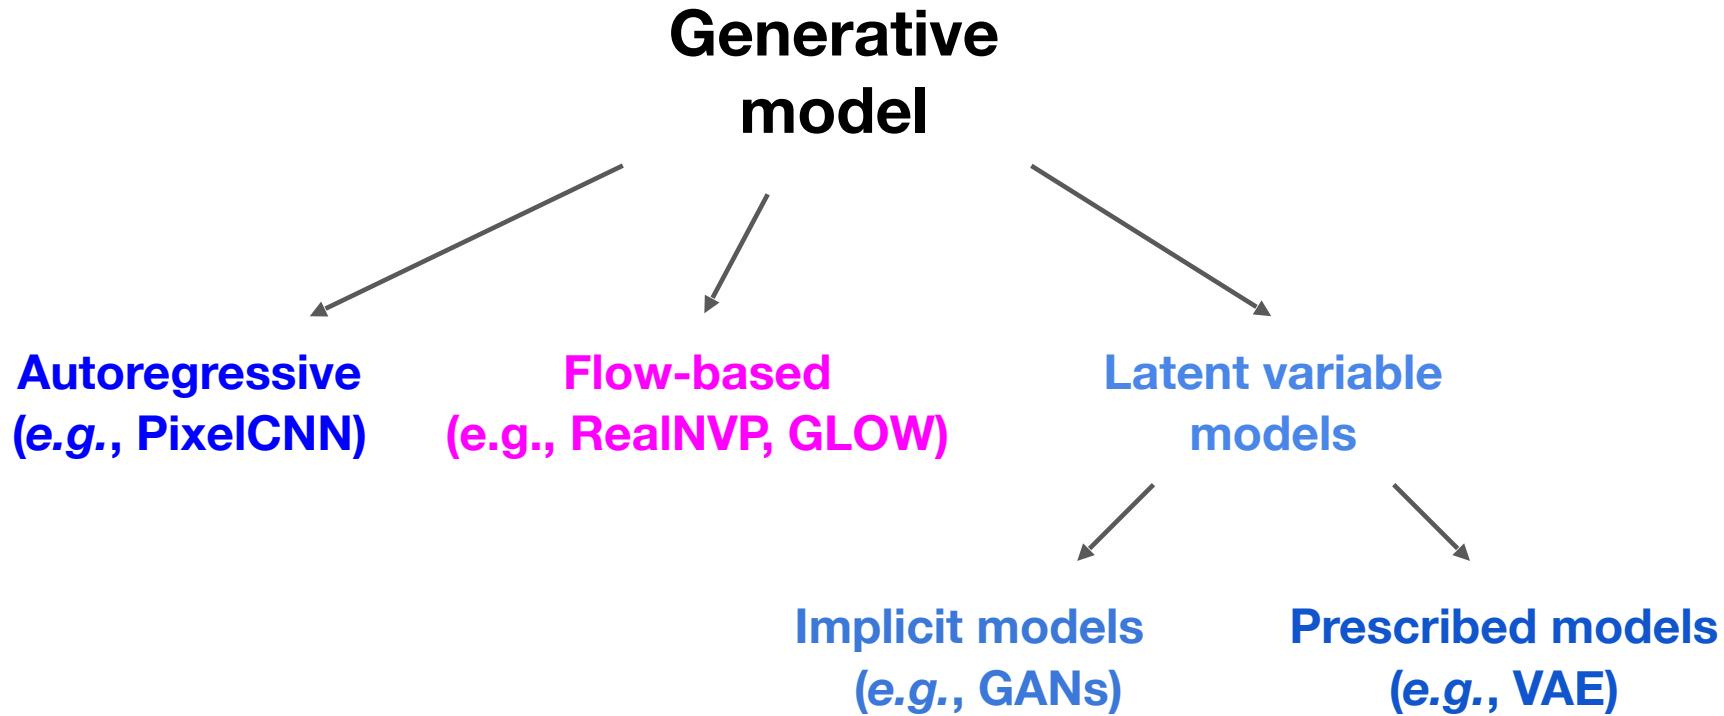


## Audio analysis



## Medical data

# HOW TO FORMULATE DEEP GENERATIVE MODELS?



# HOW TO FORMULATE DEEP GENERATIVE MODELS?

	Training	Likelihood	Sampling	Compression
Autoregressive models (e.g., PixelCNN)	Stable	Exact	Slow	No
Flow-based models (e.g., RealNVP)	Stable	Exact	Fast/Slow	No
Implicit models (e.g., GANs)	Unstable	No	Fast	No
Prescribed models (e.g., VAEs)	Stable	Approximate	Fast	Yes

# DEEP LATENT VARIABLE MODELS

# GENERATIVE MODELING IN HIGH-DIM

Modeling in high-dimensional spaces is difficult.



# GENERATIVE MODELING IN HIGH-DIM

Modeling in high-dimensional spaces is difficult.



Modeling in high-dimensional spaces is difficult.

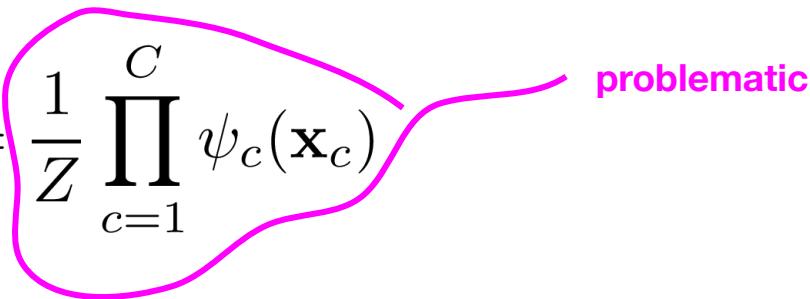
Modeling **all dependencies** among pixels:

$$p(\mathbf{x}) = \frac{1}{Z} \prod_{c=1}^C \psi_c(\mathbf{x}_c)$$

# GENERATIVE MODELING IN HIGH-DIM

Modeling in high-dimensional spaces is difficult.

Modeling **all dependencies** among pixels:

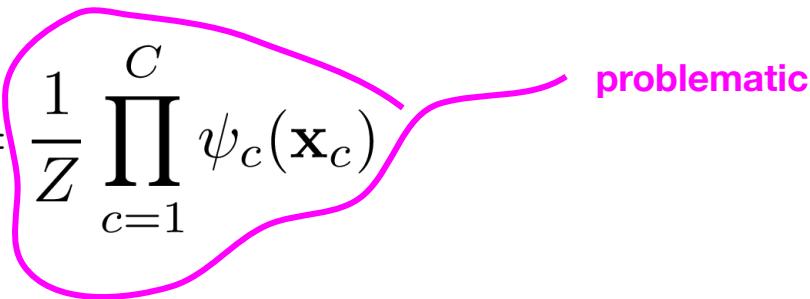
$$p(\mathbf{x}) = \frac{1}{Z} \prod_{c=1}^C \psi_c(\mathbf{x}_c)$$


problematic

# GENERATIVE MODELING IN HIGH-DIM

Modeling in high-dimensional spaces is difficult.

Modeling **all dependencies** among pixels:

$$p(\mathbf{x}) = \frac{1}{Z} \prod_{c=1}^C \psi_c(\mathbf{x}_c)$$


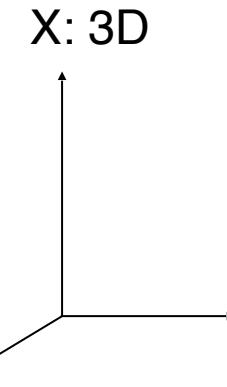
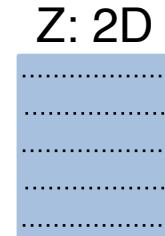
problematic

A possible **solution**: **Latent Variable Models!**

# GENERATIVE MODELING WITH LATENT VARIABLES

Generative process:

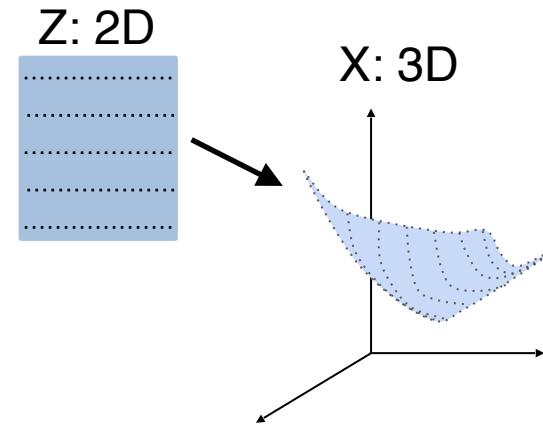
1.  $\mathbf{z} \sim p_\lambda(\mathbf{z})$
2.  $\mathbf{x} \sim p_\theta(\mathbf{x}|\mathbf{z})$



# GENERATIVE MODELING WITH LATENT VARIABLES

Generative process:

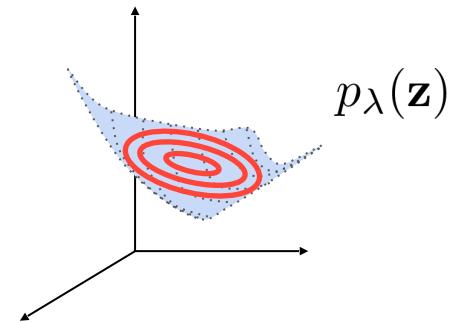
1.  $\mathbf{z} \sim p_\lambda(\mathbf{z})$
2.  $\mathbf{x} \sim p_\theta(\mathbf{x}|\mathbf{z})$



# GENERATIVE MODELING WITH LATENT VARIABLES

Generative process:

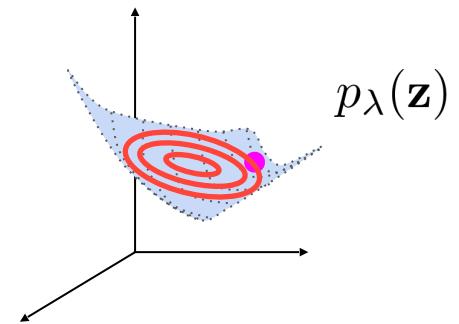
1.  $\mathbf{z} \sim p_\lambda(\mathbf{z})$
2.  $\mathbf{x} \sim p_\theta(\mathbf{x}|\mathbf{z})$



# GENERATIVE MODELING WITH LATENT VARIABLES

Generative process:

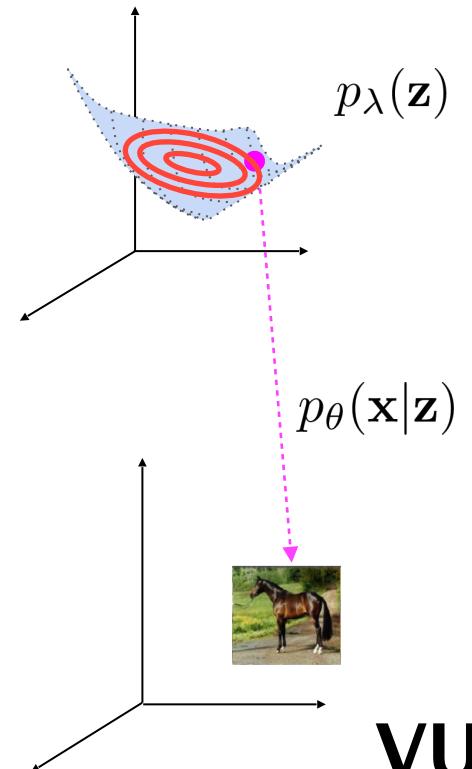
1.  $\mathbf{z} \sim p_\lambda(\mathbf{z})$
2.  $\mathbf{x} \sim p_\theta(\mathbf{x}|\mathbf{z})$



# GENERATIVE MODELING WITH LATENT VARIABLES

Generative process:

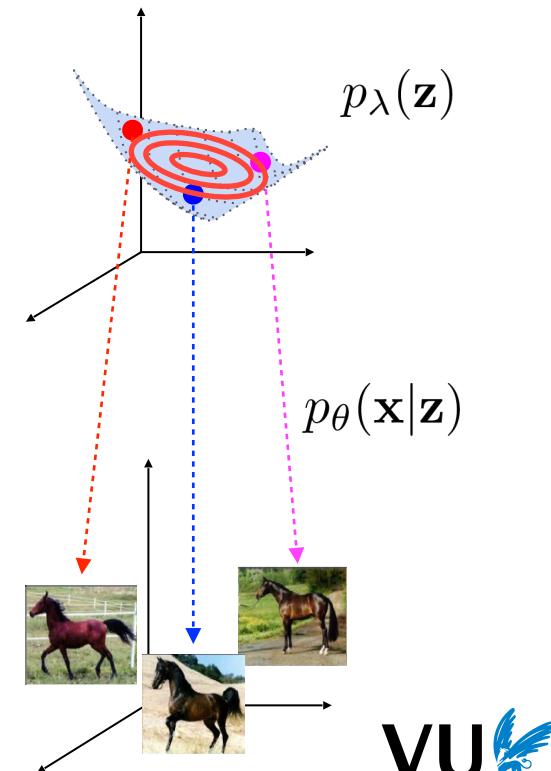
1.  $\mathbf{z} \sim p_\lambda(\mathbf{z})$
2.  $\mathbf{x} \sim p_\theta(\mathbf{x}|\mathbf{z})$



# GENERATIVE MODELING WITH LATENT VARIABLES

Generative process:

1.  $\mathbf{z} \sim p_\lambda(\mathbf{z})$
2.  $\mathbf{x} \sim p_\theta(\mathbf{x}|\mathbf{z})$



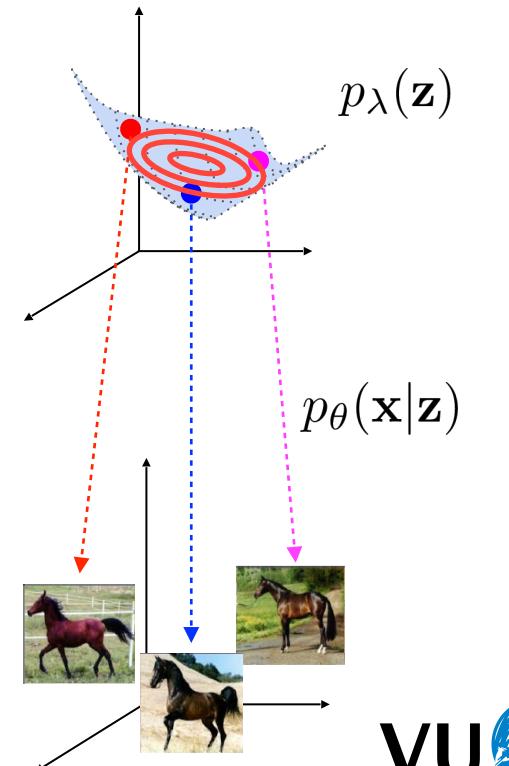
# GENERATIVE MODELING WITH LATENT VARIABLES

Generative process:

1.  $\mathbf{z} \sim p_\lambda(\mathbf{z})$
2.  $\mathbf{x} \sim p_\theta(\mathbf{x}|\mathbf{z})$

The log-likelihood function:

$$\log p_\vartheta(\mathbf{x}) = \log \int p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z}$$



# GENERATIVE MODELING WITH LATENT VARIABLES

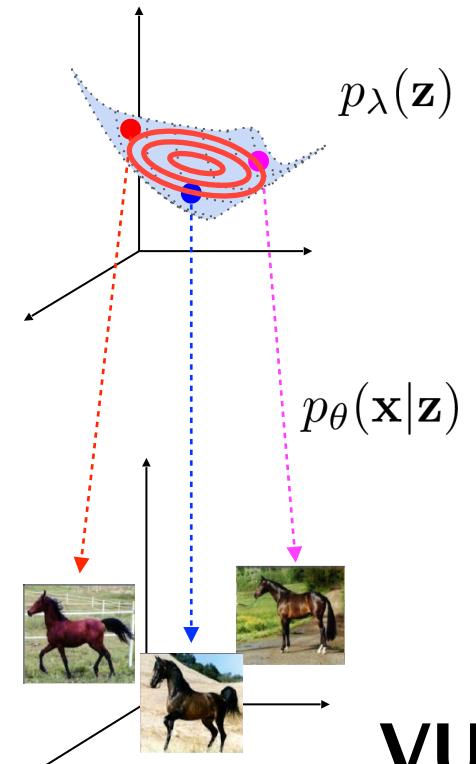
Generative process:

1.  $\mathbf{z} \sim p_\lambda(\mathbf{z})$
2.  $\mathbf{x} \sim p_\theta(\mathbf{x}|\mathbf{z})$

The log-likelihood function:

$$\log p_\vartheta(\mathbf{x}) = \log \underbrace{\int p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z}}$$

**How to train such model efficiently?**



# LINEAR LATENT VARIABLE MODELS

Let us assume:  $p(\mathbf{z}) = \mathcal{N}(0, \mathbf{I})$ .

# LINEAR LATENT VARIABLE MODELS

Let us assume:  $p(\mathbf{z}) = \mathcal{N}(0, \mathbf{I})$ .

And a linear transformation ( $\mathbf{W} \in \mathbb{R}^{D \times M}$ ):

$$\mathbf{x} = \mathbf{W}\mathbf{z} + \mu + \varepsilon, \text{ where } \varepsilon \sim \mathcal{N}(0, \sigma^2 \mathbf{I})$$

# LINEAR LATENT VARIABLE MODELS

Let us assume:  $p(\mathbf{z}) = \mathcal{N}(0, \mathbf{I})$ .

And a linear transformation ( $\mathbf{W} \in \mathbb{R}^{D \times M}$ ):

$$\mathbf{x} = \mathbf{W}\mathbf{z} + \mu + \varepsilon, \text{ where } \varepsilon \sim \mathcal{N}(0, \sigma^2 \mathbf{I})$$

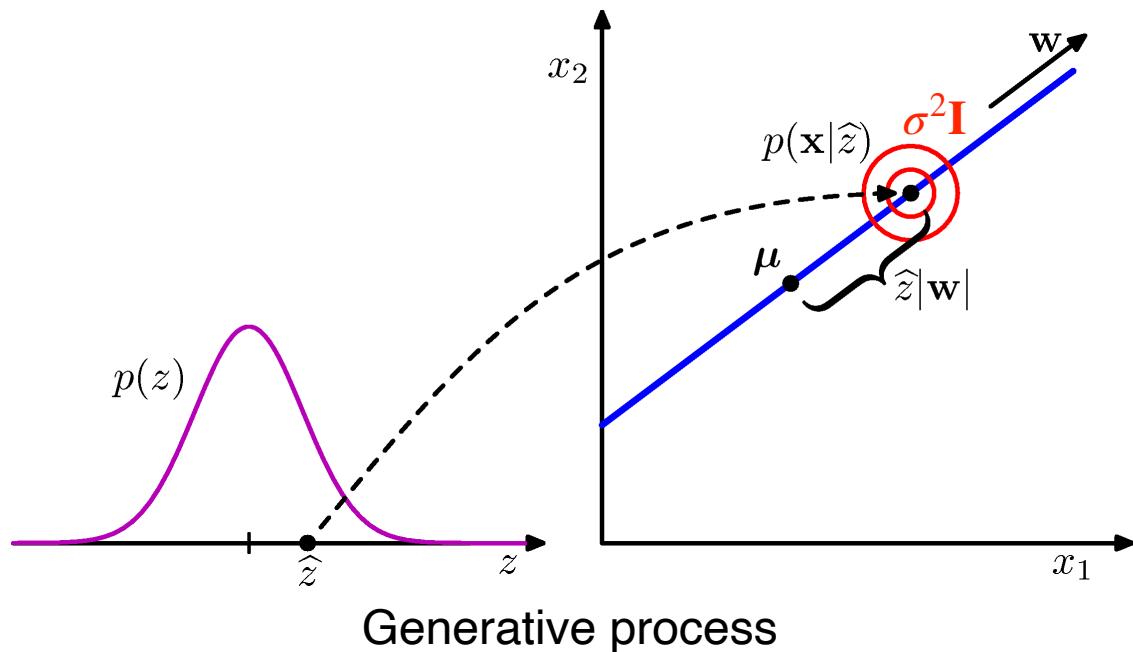
that results in the following conditional distribution:

$$p(\mathbf{x} | \mathbf{z}) = \mathcal{N}(\mathbf{W}\mathbf{z} + \mu, \sigma^2 \mathbf{I})$$

# LINEAR LATENT VARIABLE MODELS

Let us assume:  $p(\mathbf{z}) = \mathcal{N}(0, \mathbf{I})$ .

And the following conditional distribution:  $p(\mathbf{x} | \mathbf{z}) = \mathcal{N}(\mathbf{Wz} + \boldsymbol{\mu}, \sigma^2 \mathbf{I})$



Now, the question is how to calculate the log-likelihood:

$$p(\mathbf{x}) = \int p(\mathbf{x} | \mathbf{z})p(\mathbf{z})d\mathbf{z}$$

# LINEAR LATENT VARIABLE MODELS

Now, the question is how to calculate the log-likelihood:

$$p(\mathbf{x}) = \int \underbrace{p(\mathbf{x} | \mathbf{z})}_{\text{Gaussian}} \underbrace{p(\mathbf{z})}_{\text{Gaussian}} d\mathbf{z}$$

# LINEAR LATENT VARIABLE MODELS

Now, the question is how to calculate the log-likelihood:

$$p(\mathbf{x}) = \int \underbrace{p(\mathbf{x} | \mathbf{z})}_{\text{Gaussian}} \underbrace{p(\mathbf{z})}_{\text{Gaussian}} d\mathbf{z}$$

$$= \mathcal{N}(\mu, \mathbf{W}\mathbf{W}^T + \sigma^2\mathbf{I})$$

## Marginal and Conditional Gaussians

Given a marginal Gaussian distribution for  $\mathbf{x}$  and a conditional Gaussian distribution for  $\mathbf{y}$  given  $\mathbf{x}$  in the form

$$p(\mathbf{x}) = \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}, \boldsymbol{\Lambda}^{-1}) \quad (2.113)$$

$$p(\mathbf{y} | \mathbf{x}) = \mathcal{N}(\mathbf{y} | \mathbf{A}\mathbf{x} + \mathbf{b}, \mathbf{L}^{-1}) \quad (2.114)$$

the marginal distribution of  $\mathbf{y}$  and the conditional distribution of  $\mathbf{x}$  given  $\mathbf{y}$  are given by

$$p(\mathbf{y}) = \mathcal{N}(\mathbf{y} | \mathbf{A}\boldsymbol{\mu} + \mathbf{b}, \mathbf{L}^{-1} + \mathbf{A}\boldsymbol{\Lambda}^{-1}\mathbf{A}^T) \quad (2.115)$$

$$p(\mathbf{x} | \mathbf{y}) = \mathcal{N}(\mathbf{x} | \boldsymbol{\Sigma}\{\mathbf{A}^T\mathbf{L}(\mathbf{y} - \mathbf{b}) + \boldsymbol{\Lambda}\boldsymbol{\mu}\}, \boldsymbol{\Sigma}) \quad (2.116)$$

where

$$\boldsymbol{\Sigma} = (\boldsymbol{\Lambda} + \mathbf{A}^T\mathbf{L}\mathbf{A})^{-1}. \quad (2.117)$$

# LINEAR LATENT VARIABLE MODELS

Now, the question is how to calculate the log-likelihood:

$$p(\mathbf{x}) = \int \underbrace{p(\mathbf{x} | \mathbf{z})}_{\text{Gaussian}} \underbrace{p(\mathbf{z})}_{\text{Gaussian}} d\mathbf{z}$$

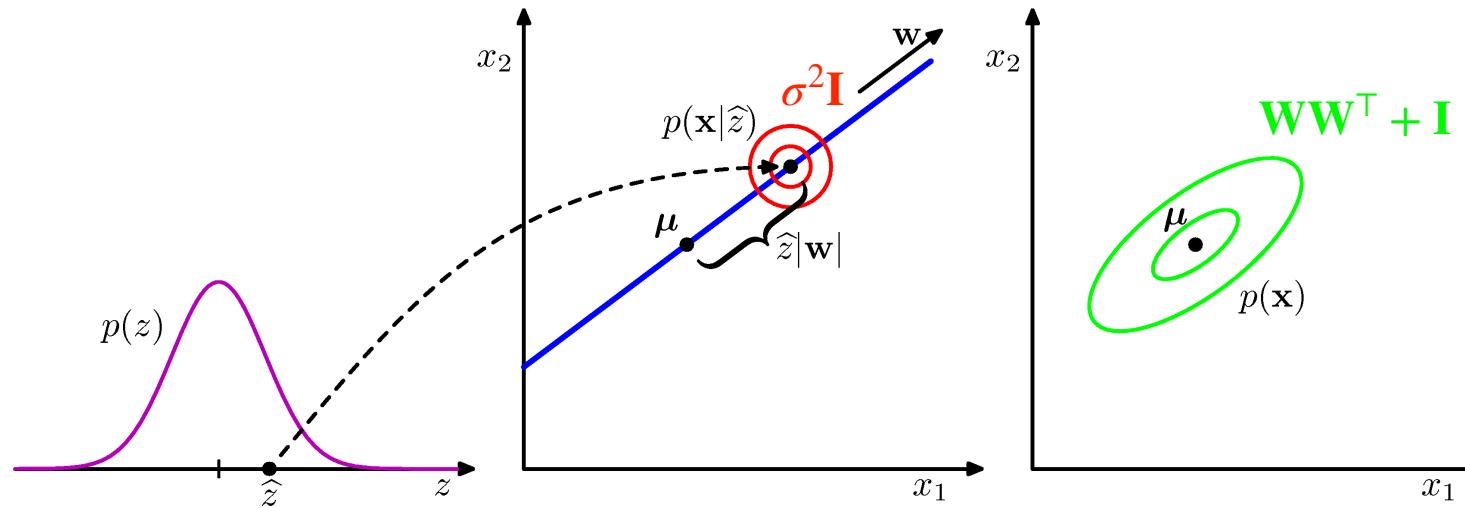
$$= \mathcal{N}(\mu, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I})$$

The integral is tractable, and  
it is again Gaussian!

# LINEAR LATENT VARIABLE MODELS

Now, the question is how to calculate the log-likelihood:

$$\begin{aligned} p(\mathbf{x}) &= \int p(\mathbf{x} | \mathbf{z}) p(\mathbf{z}) d\mathbf{z} \\ &= \mathcal{N}(\mu, \mathbf{W}\mathbf{W}^T + \sigma^2 \mathbf{I}) \end{aligned}$$



Since the model is linear, and all distributions are Gaussians, we can also calculate the posterior over  $\mathbf{z}$ :

$$p(\mathbf{z} \mid \mathbf{x}) = \mathcal{N}\left(\mathbf{M}^{-1}\mathbf{W}^\top(\mathbf{x} - \boldsymbol{\mu}), \sigma^{-2}\mathbf{M}\right)$$

where:

$$\mathbf{M} = \mathbf{W}^\top\mathbf{W} + \sigma^2\mathbf{I}$$

# LINEAR LATENT VARIABLE MODELS

Since the model is linear, and all distributions are Gaussians, we can also calculate the posterior over  $\mathbf{z}$ :

$$p(\mathbf{z} | \mathbf{x}) = \mathcal{N}\left(\mathbf{M}^{-1}\mathbf{W}^\top(\mathbf{x} - \boldsymbol{\mu}), \sigma^{-2}\mathbf{M}\right)$$

where:

$$\mathbf{M} = \mathbf{W}^\top\mathbf{W} + \sigma^2\mathbf{I}$$

## Marginal and Conditional Gaussians

Given a marginal Gaussian distribution for  $\mathbf{x}$  and a conditional Gaussian distribution for  $\mathbf{y}$  given  $\mathbf{x}$  in the form

$$p(\mathbf{x}) = \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}, \boldsymbol{\Lambda}^{-1}) \quad (2.113)$$

$$p(\mathbf{y} | \mathbf{x}) = \mathcal{N}(\mathbf{y} | \mathbf{A}\mathbf{x} + \mathbf{b}, \mathbf{L}^{-1}) \quad (2.114)$$

the marginal distribution of  $\mathbf{y}$  and the conditional distribution of  $\mathbf{x}$  given  $\mathbf{y}$  are given by

$$p(\mathbf{y}) = \mathcal{N}(\mathbf{y} | \mathbf{A}\boldsymbol{\mu} + \mathbf{b}, \mathbf{L}^{-1} + \mathbf{A}\boldsymbol{\Lambda}^{-1}\mathbf{A}^\top) \quad (2.115)$$

$$p(\mathbf{x} | \mathbf{y}) = \mathcal{N}(\mathbf{x} | \boldsymbol{\Sigma}\{\mathbf{A}^\top\mathbf{L}(\mathbf{y} - \mathbf{b}) + \boldsymbol{\Lambda}\boldsymbol{\mu}\}, \boldsymbol{\Sigma}) \quad (2.116)$$

$$\boldsymbol{\Sigma} = (\boldsymbol{\Lambda} + \mathbf{A}^\top\mathbf{L}\mathbf{A})^{-1}. \quad (2.117)$$

# PPCA: PROBABILISTIC PRINCIPAL COMPONENT ANALYSIS

The final model is the following ( $\mathbf{W} \in \mathbb{R}^{D \times M}$ ):

$$p(\mathbf{z}) = \mathcal{N}(0, \mathbf{I})$$

$$p(\mathbf{x} | \mathbf{z}) = \mathcal{N}(\mathbf{Wz} + \boldsymbol{\mu}, \sigma^2 \mathbf{I})$$

$$p(\mathbf{z} | \mathbf{x}) = \mathcal{N}\left(\mathbf{M}^{-1} \mathbf{W}^\top (\mathbf{x} - \boldsymbol{\mu}), \sigma^{-2} \mathbf{M}\right)$$

where  $\mathbf{M} = \mathbf{W}^\top \mathbf{W} + \sigma^2 \mathbf{I}$ .

and the marginal distribution:

$$p(\mathbf{x}) = \mathcal{N}(\boldsymbol{\mu}, \mathbf{WW}^\top + \sigma^2 \mathbf{I})$$

## LEARNING PPCA: THE LIKELIHOOD FUNCTION (IID DATA)

The logarithm of the likelihood function:

$$\ln p(\mathbf{X} | \mu, \mathbf{W}, \sigma^2) = \sum_{n=1}^N \ln p(\mathbf{x}_n | \mathbf{W}, \mu, \sigma^2)$$

## LEARNING PPCA: THE LIKELIHOOD FUNCTION (IID DATA)

The logarithm of the likelihood function:

$$\ln p(\mathbf{X} | \mu, \mathbf{W}, \sigma^2) = \sum_{n=1}^N \ln p(\mathbf{x}_n | \mathbf{W}, \mu, \sigma^2)$$

**REMEMBER: Everything is Gaussian!**

## LEARNING PPCA: THE LIKELIHOOD FUNCTION (IID DATA)

The logarithm of the likelihood function:

$$\begin{aligned}\ln p(\mathbf{X} | \mu, \mathbf{W}, \sigma^2) &= \sum_{n=1}^N \ln p(\mathbf{x}_n | \mathbf{W}, \mu, \sigma^2) \\ &= -\frac{ND}{2} \ln(2\pi) - \frac{N}{2} \ln |\mathbf{C}| - \frac{1}{2} \sum_{n=1}^N (\mathbf{x}_n - \mu)^\top \mathbf{C}^{-1} (\mathbf{x}_n - \mu)\end{aligned}$$

where

$$\mathbf{C} = \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}$$

$$\mathbf{C}^{-1} = \sigma^{-2}\mathbf{I} - \sigma^{-2}\mathbf{W}\mathbf{M}^{-1}\mathbf{W}^\top$$

## LEARNING PPCA: THE LIKELIHOOD FUNCTION (IID DATA)

The logarithm of the likelihood function:

$$\begin{aligned}\ln p(\mathbf{X} | \mu, \mathbf{W}, \sigma^2) &= \sum_{n=1}^N \ln p(\mathbf{x}_n | \mathbf{W}, \mu, \sigma^2) \\ &= -\frac{ND}{2} \ln(2\pi) - \frac{N}{2} \ln |\mathbf{C}| - \frac{1}{2} \sum_{n=1}^N (\mathbf{x}_n - \mu)^\top \mathbf{C}^{-1} (\mathbf{x}_n - \mu)\end{aligned}$$

where

$$\mathbf{C} = \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}$$

$$\mathbf{C}^{-1} = \underbrace{\sigma^{-1}\mathbf{I}}_{\text{Inverting C (DxD)}} - \underbrace{\sigma^{-2}\mathbf{W}\mathbf{M}^{-1}\mathbf{W}^\top}_{\text{reduces to inverting M (MxM)}}.$$

Inverting  $\mathbf{C}$  ( $D \times D$ ) reduces to inverting  $\mathbf{M}$  ( $M \times M$ ).

## LEARNING PPCA: THE LIKELIHOOD FUNCTION (IID DATA)

It is possible to calculate the solution analytically:

$$\mathbf{W}_{\text{ML}} = \mathbf{U}_M (\mathbf{L}_M - \sigma^2 \mathbf{I})^{1/2} \mathbf{R}$$

$$\sigma_{\text{ML}}^2 = \frac{1}{D - M} \sum_{i=M+1}^D \lambda_i$$

where:

$\mathbf{U}_M$  - is a  $D \times M$  matrix whose columns are eigenvectors of  $\mathbf{S}$

$\mathbf{L}_M$  - is a  $M \times M$  diagonal matrix whose elements are eigenvalues of  $\mathbf{S}$

$\mathbf{S} = \frac{1}{N} \sum_{n=1}^N (\mathbf{x}_n - \bar{\mathbf{x}}) (\mathbf{x}_n - \bar{\mathbf{x}})^T$  - is the sample covariance matrix

# LEARNING PPCA: THE LIKELIHOOD FUNCTION (IID DATA)

It is possible to calculate the solution analytically:

$$\mathbf{W}_{\text{ML}} = \mathbf{U}_M (\mathbf{L}_M - \sigma^2 \mathbf{I})^{1/2} \underbrace{\mathbf{R}}_{\text{any orthogonal matrix}}$$

$$\sigma_{\text{ML}}^2 = \underbrace{\frac{1}{D-M} \sum_{i=M+1}^D \lambda_i}_{\text{The average variance of discarded dimensions.}}$$

where:

$\mathbf{U}_M$  - is a  $D \times M$  matrix whose columns are eigenvectors of  $\mathbf{S}$

$\mathbf{L}_M$  - is a  $M \times M$  diagonal matrix whose elements are eigenvalues of  $\mathbf{S}$

$\mathbf{S} = \frac{1}{N} \sum_{n=1}^N (\mathbf{x}_n - \bar{\mathbf{x}}) (\mathbf{x}_n - \bar{\mathbf{x}})^T$  - is the sample covariance matrix

## LEARNING PPCA: THE LIKELIHOOD FUNCTION (IID DATA)

- If we use the **eigendecomposition** of the sample covariance matrix, then we can simply take  $\mathbf{R} = \mathbf{I}$ .

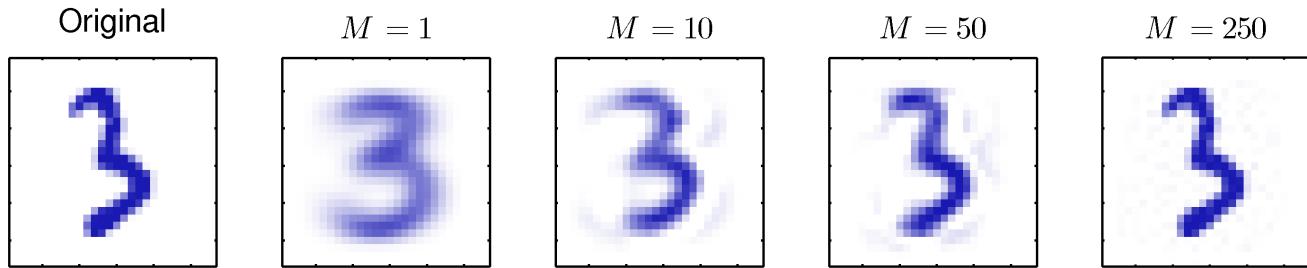
## LEARNING PPCA: THE LIKELIHOOD FUNCTION (IID DATA)

- If we use the **eigendecomposition** of the sample covariance matrix, then we can simply take  $\mathbf{R} = \mathbf{I}$ .
- In practice, the complexity of the eigendecomposition is  $O(D^3)$ , and the complexity of calculating the covariance matrix is  $O(ND^2)$ .

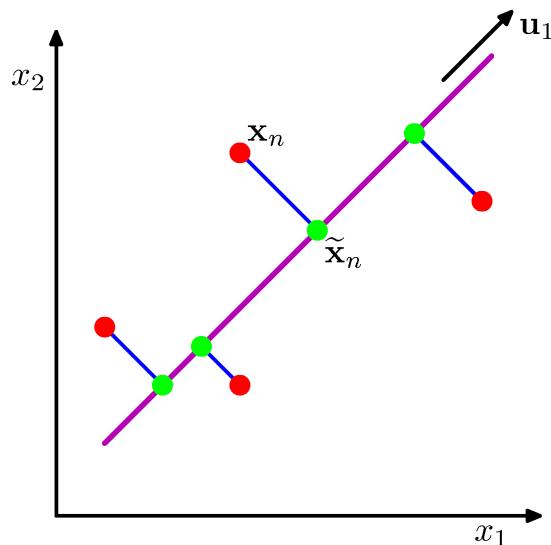
- If we use the **eigendecomposition** of the sample covariance matrix, then we can simply take  $\mathbf{R} = \mathbf{I}$ .
- In practice, the complexity of the eigendecomposition is  $O(D^3)$ , and the complexity of calculating the covariance matrix is  $O(ND^2)$ .
- If we have large problems (i.e.,  $D > 1000, N > 1000$ ), then we can use:
  - ▶ **Expectation-Maximization** (EM)
  - ▶ **Gradient-based optimization** (e.g., SGD).
- For numerical algorithms,  $\mathbf{R}$  could be arbitrary, so **no unique solution**.

# PROBABILISTIC PCA

Reconstruction:



Projection of 2D data  
onto 1D space:



## Advantages

- ✓ Exact likelihood.
- ✓ Analytical solution.
- ✓ Posterior over  $\mathbf{z}$  is analytical.
- ✓ Allows compression.
- ✓ Allows to generate.

## Disadvantages

- Linear transformation drastically limits the applicability.
- For more complex data, pPCA requires  $M$  close to  $D$  to work well.
- No analytical solution for binary data.

# VARIATIONAL AUTO-ENCODERS

# GENERATIVE MODELS

	Training	Likelihood	Sampling	Compression
Autoregressive models (e.g., PixelCNN)	Stable	Exact	Slow	No
Flow-based models (e.g., RealNVP)	Stable	Exact	Fast/Slow	No
Implicit models (e.g., GANs)	Unstable	No	Fast	No
Prescribed models (e.g., VAEs)	Stable	Approximate	Fast	Yes

# GENERATIVE MODELING WITH LATENT VARIABLES

Generative process:

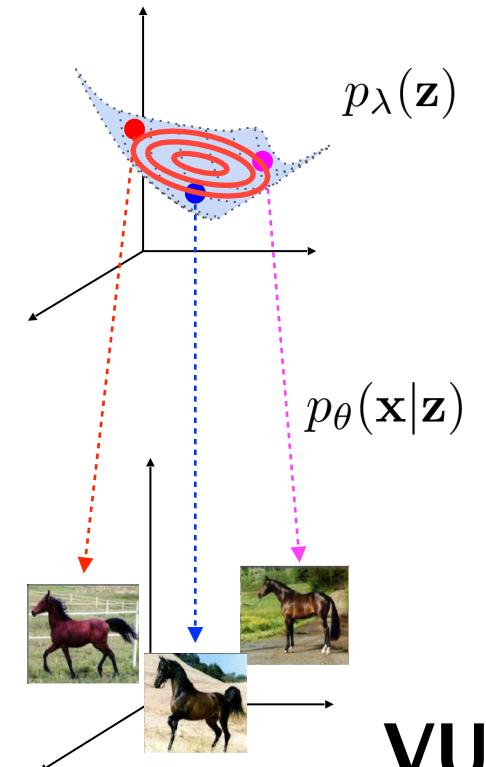
1.  $\mathbf{z} \sim p_\lambda(\mathbf{z})$
2.  $\mathbf{x} \sim p_\theta(\mathbf{x}|\mathbf{z})$

The log-likelihood function:

$$\log p_\theta(\mathbf{x}) = \log \underbrace{\int p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z}}$$

**How to train such model efficiently?**

**Now we consider non-linear transformations.**



# LATENT-VARIABLE MODELS WITH NON-LINEAR TRANSFORMATIONS

Let us assume:  $p(\mathbf{z}) = \mathcal{N}(0, \mathbf{I})$ .

For the pPCA:  $p(\mathbf{x} | \mathbf{z}) = \mathcal{N}(\mathbf{Wz} + \boldsymbol{\mu}, \sigma^2 \mathbf{I})$

Let us assume:  $p(\mathbf{z}) = \mathcal{N}(0, \mathbf{I})$ .

For the pPCA:  $p(\mathbf{x} | \mathbf{z}) = \mathcal{N}(\mathbf{W}\mathbf{z} + \boldsymbol{\mu}, \sigma^2 \mathbf{I})$

Now, we consider:  $p(\mathbf{x} | \mathbf{z}) = \mathcal{N}(f(\mathbf{z}; \mathbf{W}), \sigma^2 \mathbf{I})$ .

Let us assume:  $p(\mathbf{z}) = \mathcal{N}(0, \mathbf{I})$ .

For the pPCA:  $p(\mathbf{x} | \mathbf{z}) = \mathcal{N}(\mathbf{Wz} + \boldsymbol{\mu}, \sigma^2 \mathbf{I})$

Now, we consider:  $p(\mathbf{x} | \mathbf{z}) = \mathcal{N}(f(\mathbf{z}; \mathbf{W}), \sigma^2 \mathbf{I})$ .

Since  $f$  could be any non-linear transformation, Prof. Bishop cannot provide us any tricks to solve the integral:

$$p(\mathbf{x}) = \int p(\mathbf{x} | \mathbf{z})p(\mathbf{z})d\mathbf{z}$$

This is an infinite mixture of Gaussians.

Let us assume:  $p(\mathbf{z}) = \mathcal{N}(0, \mathbf{I})$ .

For the pPCA:  $p(\mathbf{x} | \mathbf{z}) = \mathcal{N}(\mathbf{W}\mathbf{z} + \boldsymbol{\mu}, \sigma^2 \mathbf{I})$

Now, we consider:  $p(\mathbf{x} | \mathbf{z}) = \mathcal{N}(f(\mathbf{z}; \mathbf{W}), \sigma^2 \mathbf{I})$ .

Since  $f$  could be any non-linear transformation, Prof. Bishop cannot provide us any tricks to solve the integral:

$$p(\mathbf{x}) = \int p(\mathbf{x} | \mathbf{z})p(\mathbf{z})d\mathbf{z}$$

This is an infinite mixture of Gaussians.

BUT we can use variational inference!  
(Chapter 10 in Bishop's book 😊)

# VARIATIONAL INFERENCE FOR LATENT VARIABLE MODELS

$$\begin{aligned}\log p_\vartheta(\mathbf{x}) &= \log \int p_\theta(\mathbf{x}|\mathbf{z}) \ p_\lambda(\mathbf{z}) d\mathbf{z} \\ &= \log \int \frac{q_\phi(\mathbf{z}|\mathbf{x})}{q_\phi(\mathbf{z}|\mathbf{x})} p_\theta(\mathbf{x}|\mathbf{z}) \ p_\lambda(\mathbf{z}) d\mathbf{z} \\ &\geq \int q_\phi(\mathbf{z}|\mathbf{x}) \log \frac{p_\theta(\mathbf{x}|\mathbf{z}) \ p_\lambda(\mathbf{z})}{q_\phi(\mathbf{z}|\mathbf{x})} d\mathbf{z} \\ &= \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})} \left[ \log p_\theta(\mathbf{x}|\mathbf{z}) \right] - \text{KL}\left(q_\phi(\mathbf{z}|\mathbf{x}) \middle\| p_\lambda(\mathbf{z})\right)\end{aligned}$$

# VARIATIONAL INFERENCE FOR LATENT VARIABLE MODELS

$$\begin{aligned}\log p_\vartheta(\mathbf{x}) &= \log \int p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z} && \text{Variational posterior} \\ &= \log \int \frac{q_\phi(\mathbf{z}|\mathbf{x})}{q_\phi(\mathbf{z}|\mathbf{x})} p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z} \\ &\geq \int q_\phi(\mathbf{z}|\mathbf{x}) \log \frac{p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z})}{q_\phi(\mathbf{z}|\mathbf{x})} d\mathbf{z} \\ &= \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})} \left[ \log p_\theta(\mathbf{x}|\mathbf{z}) \right] - \text{KL}\left(q_\phi(\mathbf{z}|\mathbf{x}) \middle\| p_\lambda(\mathbf{z})\right)\end{aligned}$$

# VARIATIONAL INFERENCE FOR LATENT VARIABLE MODELS

$$\begin{aligned}\log p_\vartheta(\mathbf{x}) &= \log \int p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z} \\ &= \log \int \frac{q_\phi(\mathbf{z}|\mathbf{x})}{q_\phi(\mathbf{z}|\mathbf{x})} p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z} \\ &\geq \int q_\phi(\mathbf{z}|\mathbf{x}) \log \frac{p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z})}{q_\phi(\mathbf{z}|\mathbf{x})} d\mathbf{z} \\ &= \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})} \left[ \log p_\theta(\mathbf{x}|\mathbf{z}) \right] - \text{KL}\left(q_\phi(\mathbf{z}|\mathbf{x}) \parallel p_\lambda(\mathbf{z})\right)\end{aligned}$$

Variational posterior

We can learn a separate  $q$  for each  $\mathbf{x}$ , but it would be too complicated.

Therefore, we use amortization.

# VARIATIONAL INFERENCE FOR LATENT VARIABLE MODELS

$$\begin{aligned}\log p_\vartheta(\mathbf{x}) &= \log \int p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z} \\ &= \log \int \frac{q_\phi(\mathbf{z}|\mathbf{x})}{q_\phi(\mathbf{z}|\mathbf{x})} p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z} \\ &\geq \int q_\phi(\mathbf{z}|\mathbf{x}) \log \frac{p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z})}{q_\phi(\mathbf{z}|\mathbf{x})} d\mathbf{z} \quad \text{Jensen's inequality} \\ &= \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})} \left[ \log p_\theta(\mathbf{x}|\mathbf{z}) \right] - \text{KL}\left(q_\phi(\mathbf{z}|\mathbf{x}) \middle\| p_\lambda(\mathbf{z})\right)\end{aligned}$$

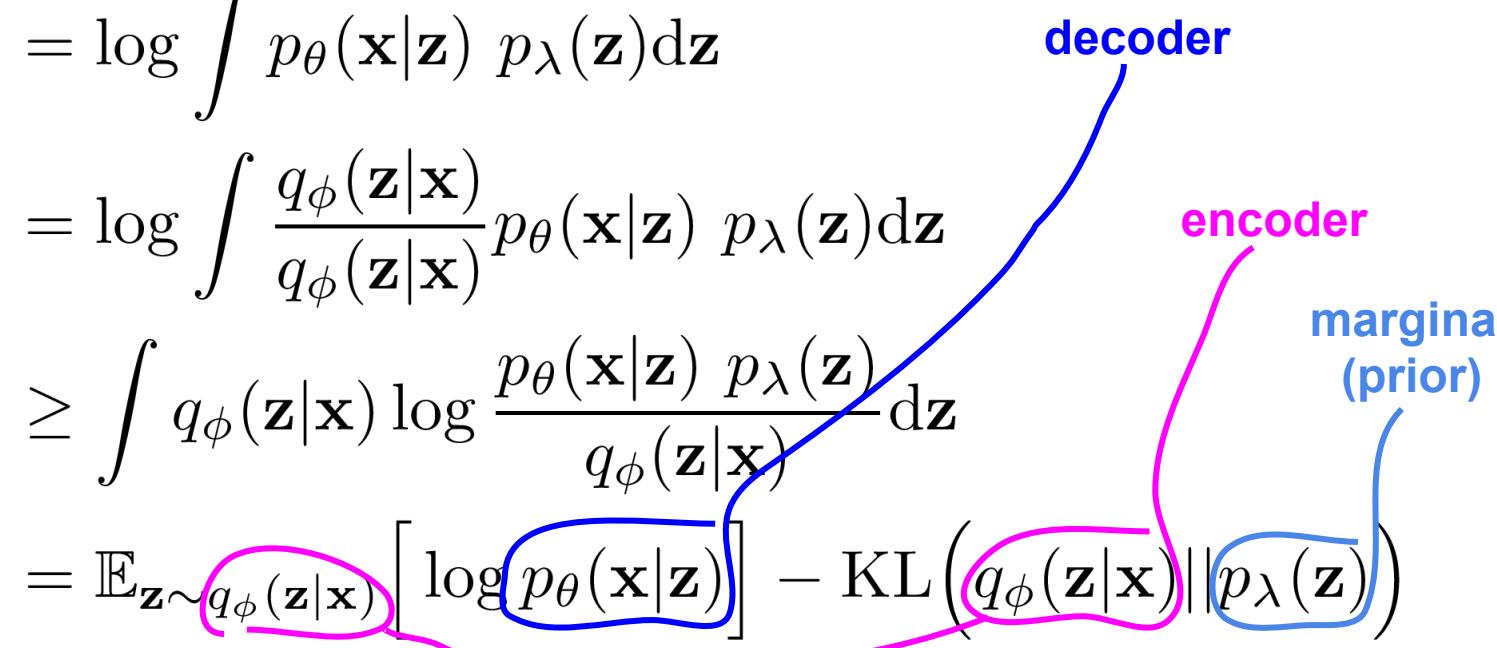
# VARIATIONAL INFERENCE FOR LATENT VARIABLE MODELS

$$\begin{aligned}\log p_\vartheta(\mathbf{x}) &= \log \int p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z} \\ &= \log \int \frac{q_\phi(\mathbf{z}|\mathbf{x})}{q_\phi(\mathbf{z}|\mathbf{x})} p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z} \\ &\geq \int q_\phi(\mathbf{z}|\mathbf{x}) \log \frac{p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z})}{q_\phi(\mathbf{z}|\mathbf{x})} d\mathbf{z} \\ &= \underbrace{\mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})} \left[ \log p_\theta(\mathbf{x}|\mathbf{z}) \right]}_{\text{Evidence Lower BOund (ELBO)}} - \text{KL}\left(q_\phi(\mathbf{z}|\mathbf{x}) \parallel p_\lambda(\mathbf{z})\right)\end{aligned}$$

# VARIATIONAL INFERENCE FOR LATENT VARIABLE MODELS

$$\begin{aligned}\log p_\vartheta(\mathbf{x}) &= \log \int p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z} \\ &= \log \int \frac{q_\phi(\mathbf{z}|\mathbf{x})}{q_\phi(\mathbf{z}|\mathbf{x})} p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z} \\ &\geq \int q_\phi(\mathbf{z}|\mathbf{x}) \log \frac{p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z})}{q_\phi(\mathbf{z}|\mathbf{x})} d\mathbf{z} \\ &= \underbrace{\mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})} \left[ \log p_\theta(\mathbf{x}|\mathbf{z}) \right]}_{\text{Reconstruction error (RE)}} - \underbrace{\text{KL}\left( q_\phi(\mathbf{z}|\mathbf{x}) \middle\| p_\lambda(\mathbf{z}) \right)}_{\text{Regularization (KL)}}\end{aligned}$$

# VARIATIONAL INFERENCE FOR LATENT VARIABLE MODELS

$$\begin{aligned}\log p_\vartheta(\mathbf{x}) &= \log \int p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z} \\ &= \log \int \frac{q_\phi(\mathbf{z}|\mathbf{x})}{q_\phi(\mathbf{z}|\mathbf{x})} p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z} \\ &\geq \int q_\phi(\mathbf{z}|\mathbf{x}) \log \frac{p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z})}{q_\phi(\mathbf{z}|\mathbf{x})} d\mathbf{z} \\ &= \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})} \left[ \log p_\theta(\mathbf{x}|\mathbf{z}) \right] - \text{KL} \left( q_\phi(\mathbf{z}|\mathbf{x}) \parallel p_\lambda(\mathbf{z}) \right)\end{aligned}$$


# VARIATIONAL INFERENCE FOR LATENT VARIABLE MODELS

$$\log p_\vartheta(\mathbf{x}) = \log \int p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z}$$

$$= \log \int \frac{q_\phi(\mathbf{z}|\mathbf{x})}{q_\phi(\mathbf{z}|\mathbf{x})} p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z}$$

$$\geq \int q_\phi(\mathbf{z}|\mathbf{x}) \log \frac{p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z})}{q_\phi(\mathbf{z}|\mathbf{x})} d\mathbf{z}$$

$$= \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})} \left[ \log p_\theta(\mathbf{x}|\mathbf{z}) \right] - \text{KL}\left(q_\phi(\mathbf{z}|\mathbf{x}) \parallel p_\lambda(\mathbf{z})\right)$$

= Variational Auto-Encoder

# VARIATIONAL INFERENCE FOR LATENT VARIABLE MODELS (A DIFFERENT DERIVATION)

$$\begin{aligned}\ln p(\mathbf{x}) &= \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} [\ln p(\mathbf{x})] \\&= \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} \left[ \ln \frac{p(\mathbf{z}|\mathbf{x})p(\mathbf{x})}{p(\mathbf{z})} \right] \\&= \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} \left[ \ln \frac{p(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p(\mathbf{z}|\mathbf{x})} \right] \\&= \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} \left[ \ln \frac{p(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p(\mathbf{z}|\mathbf{x})} \frac{q(\mathbf{z}|\mathbf{x})}{q(\mathbf{z}|\mathbf{x})} \right] \\&= \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} \left[ \ln p(\mathbf{x}|\mathbf{z}) \frac{p(\mathbf{z})}{q(\mathbf{z}|\mathbf{x})} \frac{q(\mathbf{z}|\mathbf{x})}{p(\mathbf{z}|\mathbf{x})} \right] \\&= \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} \left[ \ln p(\mathbf{x}|\mathbf{z}) - \ln \frac{q(\mathbf{z}|\mathbf{x})}{p(\mathbf{z})} + \ln \frac{q(\mathbf{z}|\mathbf{x})}{p(\mathbf{z}|\mathbf{x})} \right] \\&= \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} [\ln p(\mathbf{x}|\mathbf{z})] - KL [q(\mathbf{z}|\mathbf{x})||p(\mathbf{z})] + KL [q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x})]\end{aligned}$$

# VARIATIONAL INFERENCE FOR LATENT VARIABLE MODELS (A DIFFERENT DERIVATION)

$$\begin{aligned}\ln p(\mathbf{x}) &= \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} [\ln p(\mathbf{x})] \\&= \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} \left[ \ln \frac{p(\mathbf{z}|\mathbf{x})p(\mathbf{x})}{p(\mathbf{z})} \right] \\&= \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} \left[ \ln \frac{p(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p(\mathbf{z}|\mathbf{x})} \right] \\&= \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} \left[ \ln \frac{p(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p(\mathbf{z}|\mathbf{x})} \frac{q(\mathbf{z}|\mathbf{x})}{q(\mathbf{z}|\mathbf{x})} \right] \\&= \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} \left[ \ln p(\mathbf{x}|\mathbf{z}) \frac{p(\mathbf{z})}{q(\mathbf{z}|\mathbf{x})} \frac{q(\mathbf{z}|\mathbf{x})}{p(\mathbf{z}|\mathbf{x})} \right] \\&= \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} \left[ \ln p(\mathbf{x}|\mathbf{z}) - \ln \frac{q(\mathbf{z}|\mathbf{x})}{p(\mathbf{z})} + \ln \frac{q(\mathbf{z}|\mathbf{x})}{p(\mathbf{z}|\mathbf{x})} \right] \\&= \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} [\underbrace{\ln p(\mathbf{x}|\mathbf{z})}_{\text{ELBO}}] - KL [q(\mathbf{z}|\mathbf{x})||p(\mathbf{z})] + \underbrace{KL [q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x})]}_{\geq 0}\end{aligned}$$

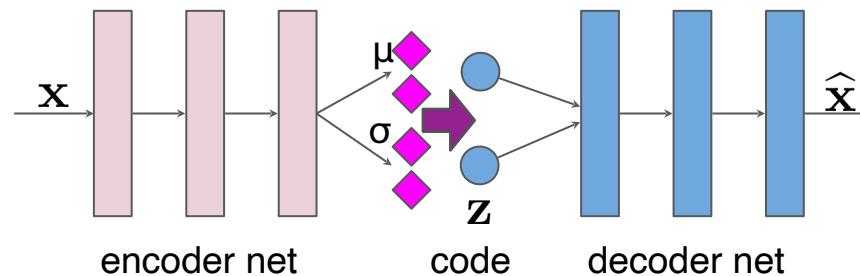
# VARIATIONAL INFERENCE FOR LATENT VARIABLE MODELS (A DIFFERENT DERIVATION)

$$\begin{aligned}\ln p(\mathbf{x}) &= \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} [\ln p(\mathbf{x})] \\&= \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} \left[ \ln \frac{p(\mathbf{z}|\mathbf{x})p(\mathbf{x})}{p(\mathbf{z})} \right] \\&= \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} \left[ \ln \frac{p(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p(\mathbf{z}|\mathbf{x})} \right] \\&= \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} \left[ \ln \frac{p(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p(\mathbf{z}|\mathbf{x})} \frac{q(\mathbf{z}|\mathbf{x})}{q(\mathbf{z}|\mathbf{x})} \right] \\&= \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} \left[ \ln p(\mathbf{x}|\mathbf{z}) \frac{p(\mathbf{z})}{q(\mathbf{z}|\mathbf{x})} \frac{q(\mathbf{z}|\mathbf{x})}{p(\mathbf{z}|\mathbf{x})} \right] \\&= \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} \left[ \ln p(\mathbf{x}|\mathbf{z}) - \ln \frac{q(\mathbf{z}|\mathbf{x})}{p(\mathbf{z})} + \ln \frac{q(\mathbf{z}|\mathbf{x})}{p(\mathbf{z}|\mathbf{x})} \right] \\&= \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} \left[ \underbrace{\ln p(\mathbf{x}|\mathbf{z})}_{\text{ELBO}} - KL [q(\mathbf{z}|\mathbf{x})||p(\mathbf{z})] + \underbrace{KL [q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x})]}_{\geq 0} \right]\end{aligned}$$

If variational posterior is poorly chosen, then the lower bound is very loose.

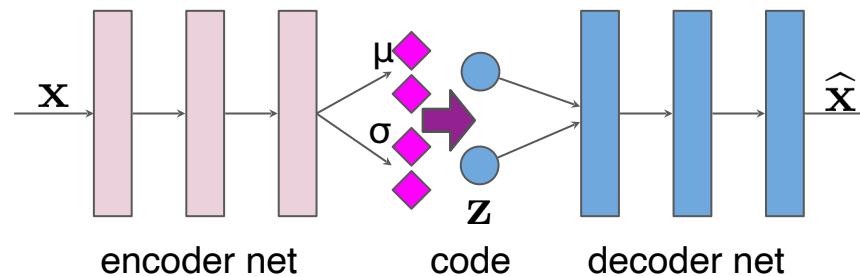
# VARIATIONAL AUTO-ENCODERS

Variational posterior (**encoder**) and the likelihood function (**decoder**) are parameterized by neural networks.



# VARIATIONAL AUTO-ENCODERS

Variational posterior (**encoder**) and the likelihood function (**decoder**) are parameterized by neural networks.



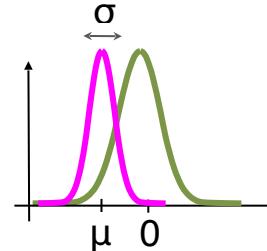
***Reparameterization trick:***

move the stochasticity to independent random variables

$$z = f(\theta, \varepsilon), \quad \varepsilon \sim p(\varepsilon)$$

e.g.  $z = \mu + \sigma \cdot \varepsilon, \quad \varepsilon \sim \mathcal{N}(0,1)$

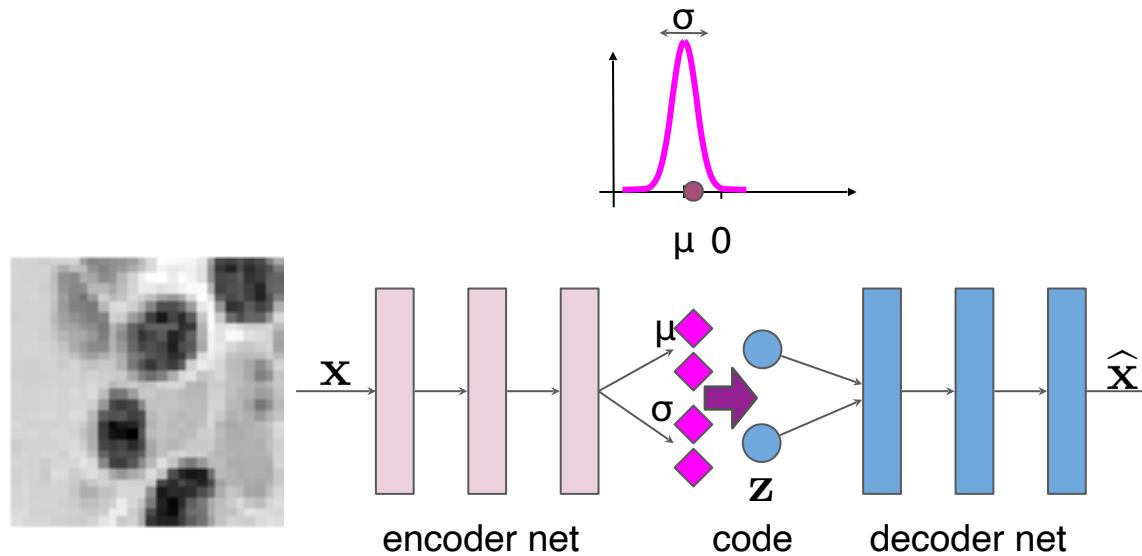
77



# VARIATIONAL AUTO-ENCODERS

VAE copies input to output through a **bottleneck**.

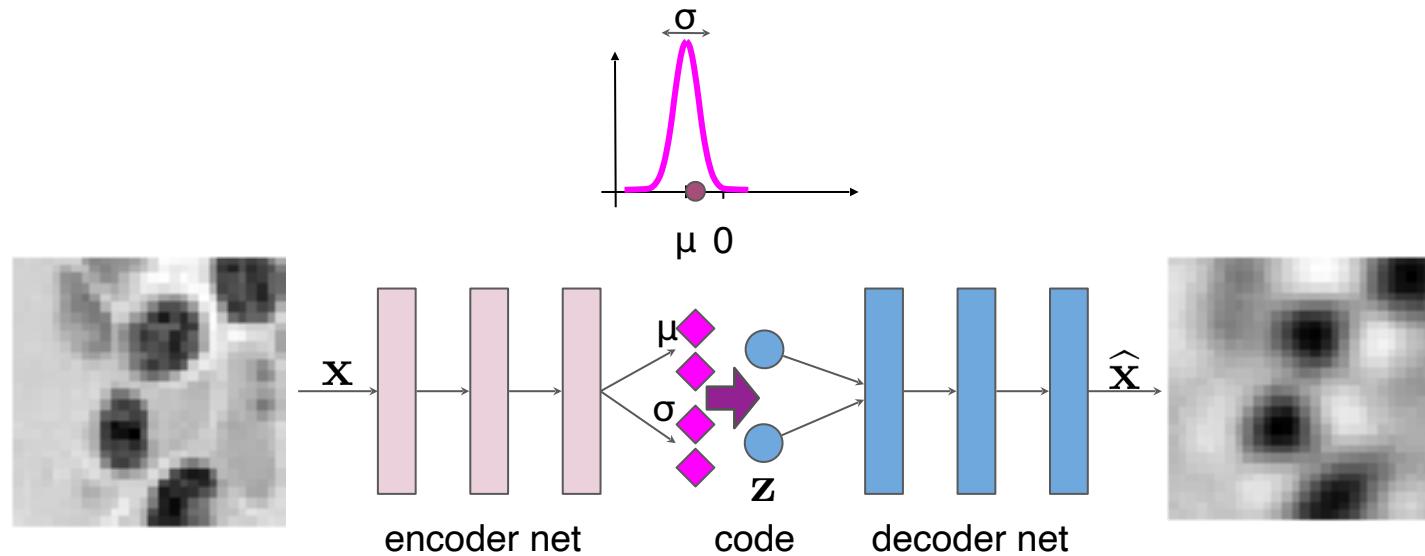
VAE learns a **code** of the data.



# VARIATIONAL AUTO-ENCODERS

VAE copies input to output through a **bottleneck**.

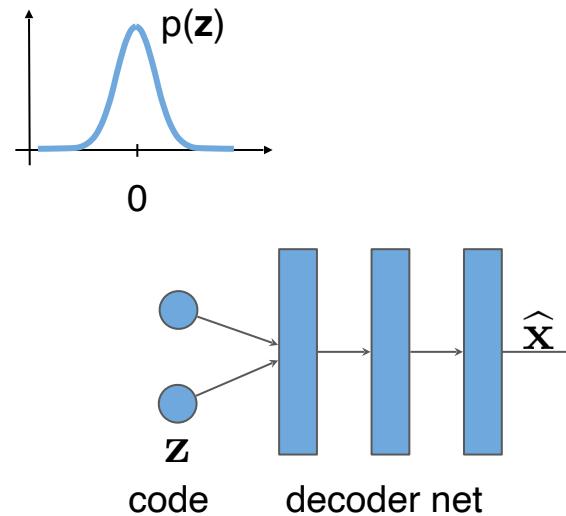
VAE learns a **code** of the data.



# VARIATIONAL AUTO-ENCODERS

VAE has a **marginal** on the latent code.

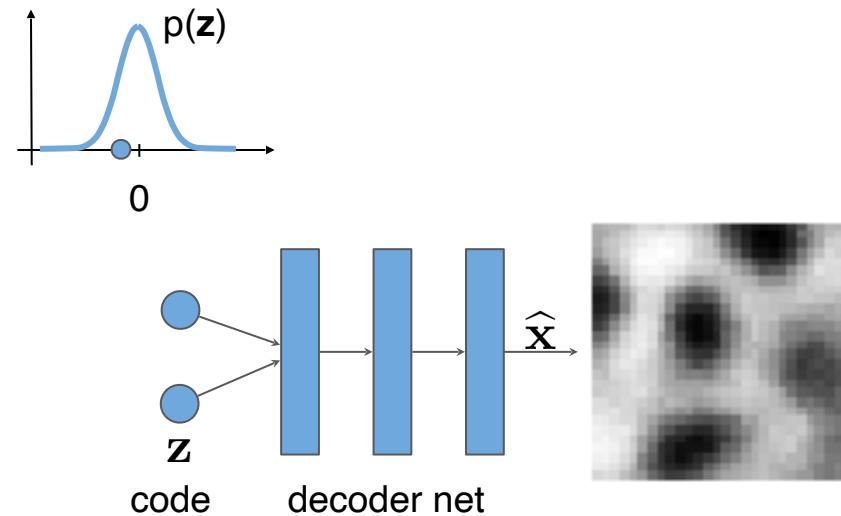
VAE can **generate** new data.



# VARIATIONAL AUTO-ENCODERS

VAE has a **marginal** on the latent code.

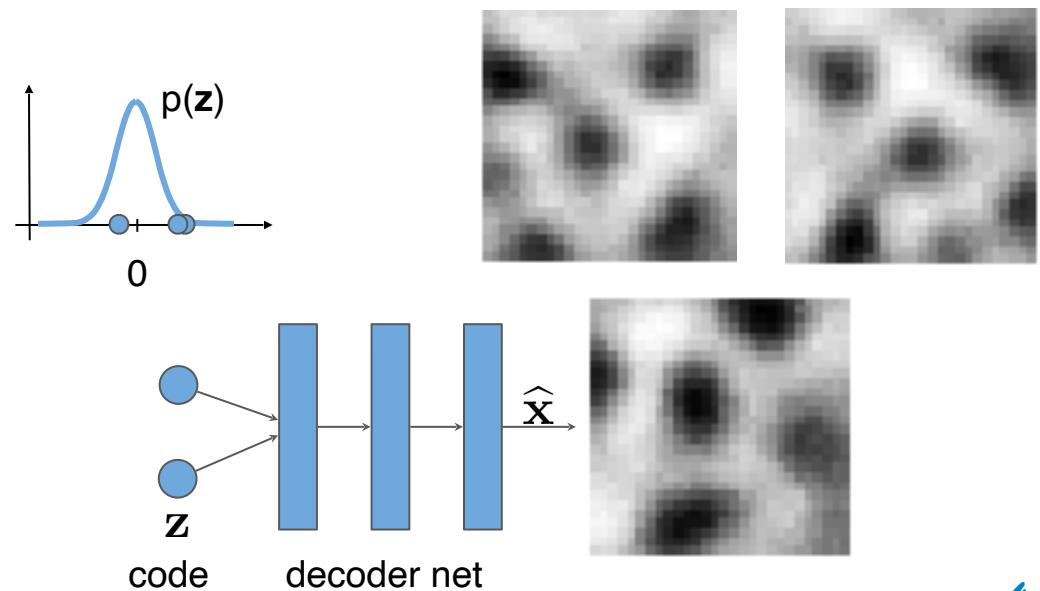
VAE can **generate** new data.



# VARIATIONAL AUTO-ENCODERS

VAE has a **marginal** on the latent code.

VAE can **generate** new data.

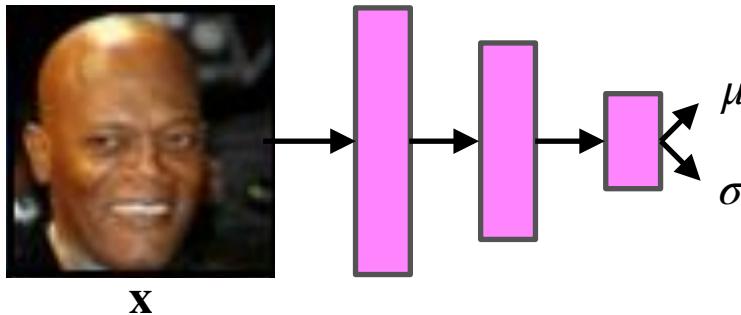


# VANILLA VAE



**X**

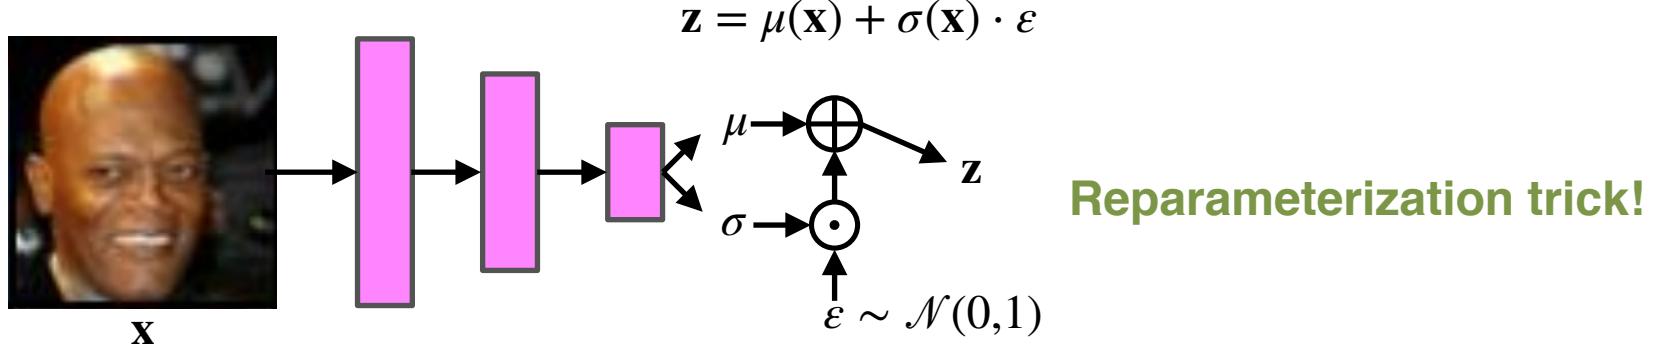
# VANILLA VAE



Example architecture for the encoder:

$\mathbf{x} \rightarrow \text{Linear}(D, 300) \rightarrow \text{ReLU} \rightarrow \text{Linear}(300, 2M) \rightarrow \text{split to 2 vectors}$

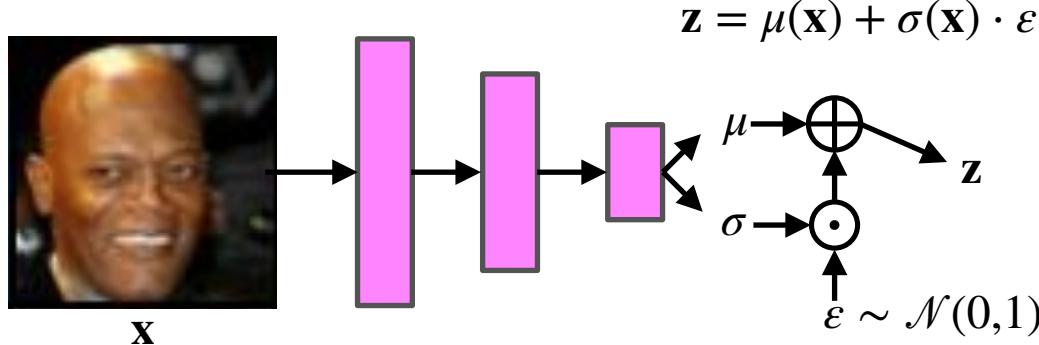
# VANILLA VAE



Example architecture for the encoder:

$\mathbf{x} \rightarrow \text{Linear}(D, 300) \rightarrow \text{ReLU} \rightarrow \text{Linear}(300, 2M) \rightarrow \text{split to 2 vectors}$

# VANILLA VAE



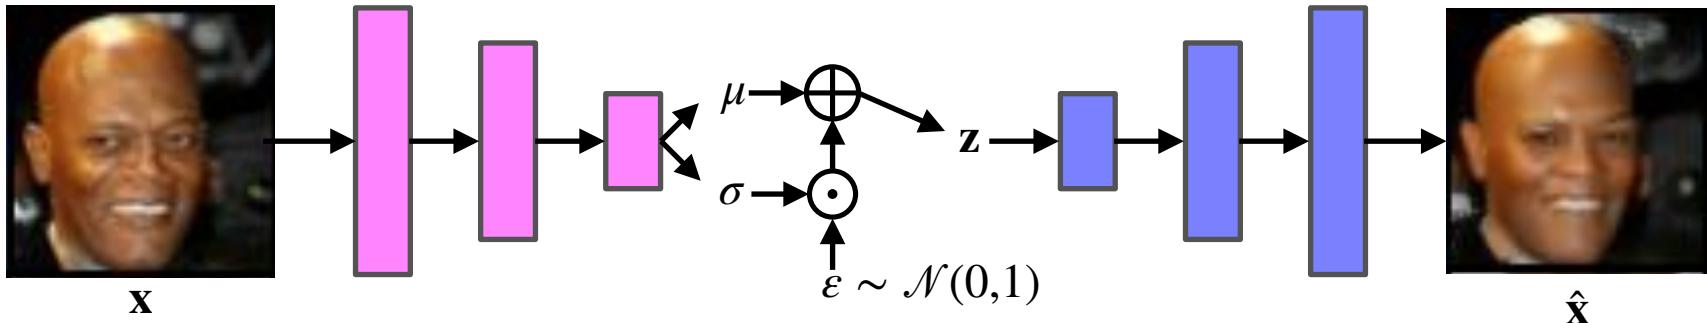
Example architecture for the encoder:

$\mathbf{x} \rightarrow \text{Linear}(D, 300) \rightarrow \text{ReLU} \rightarrow \text{Linear}(300, 2M) \rightarrow \text{split to 2 vectors}$



No non-linearity here!  
We model means and log-std  
for Gaussian.

# VANILLA VAE



Example architecture for the encoder:

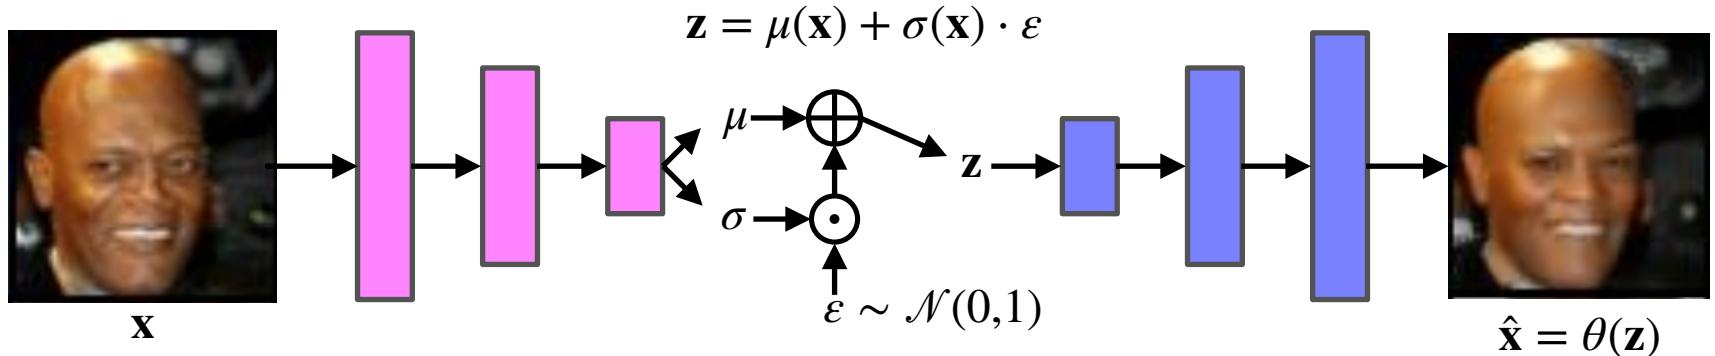
$x \rightarrow \text{Linear}(D, 300) \rightarrow \text{ReLU} \rightarrow \text{Linear}(300, 2M) \rightarrow \text{split to 2 vectors}$

Example architecture for the decoder:

$z \rightarrow \text{Linear}(M, 300) \rightarrow \text{ReLU} \rightarrow \text{Linear}(300, D) \rightarrow \text{means}$

No non-linearity here!  
We model means only.

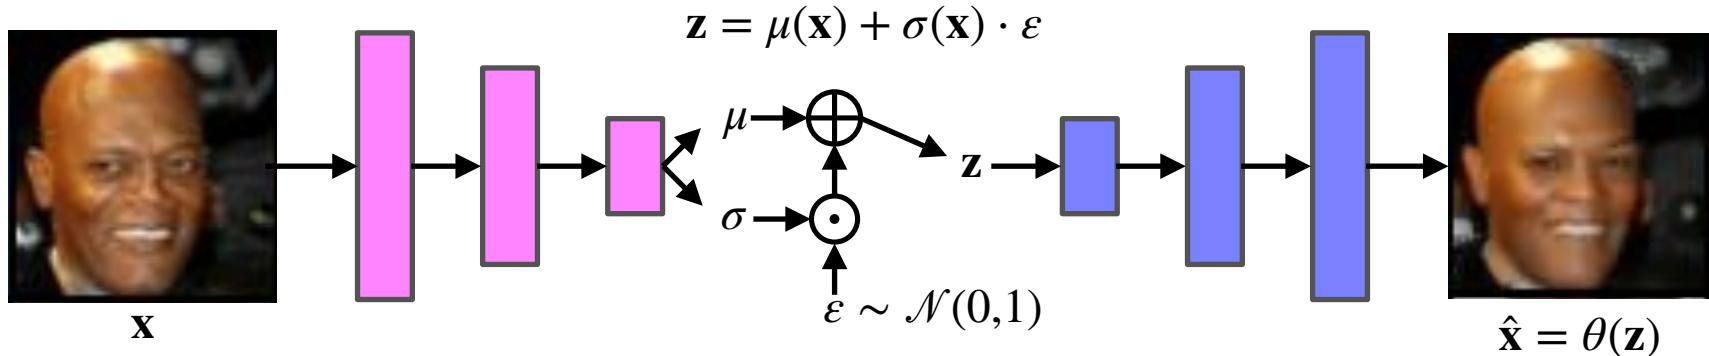
# VANILLA VAE



We approximate expected values using a single sample:

$$ELBO = \underbrace{\ln \mathcal{N}(\mathbf{x} | \theta(\mathbf{z}), 1)}_{p_\theta(\mathbf{x}|\mathbf{z})} - \left[ \underbrace{\ln \mathcal{N}(\mathbf{z} | \mu(\mathbf{x}), \sigma^2(\mathbf{x}))}_{q_\phi(\mathbf{z}|\mathbf{x})} - \underbrace{\ln \mathcal{N}(\mathbf{z} | 0, 1)}_{p_\lambda(\mathbf{z})} \right]$$

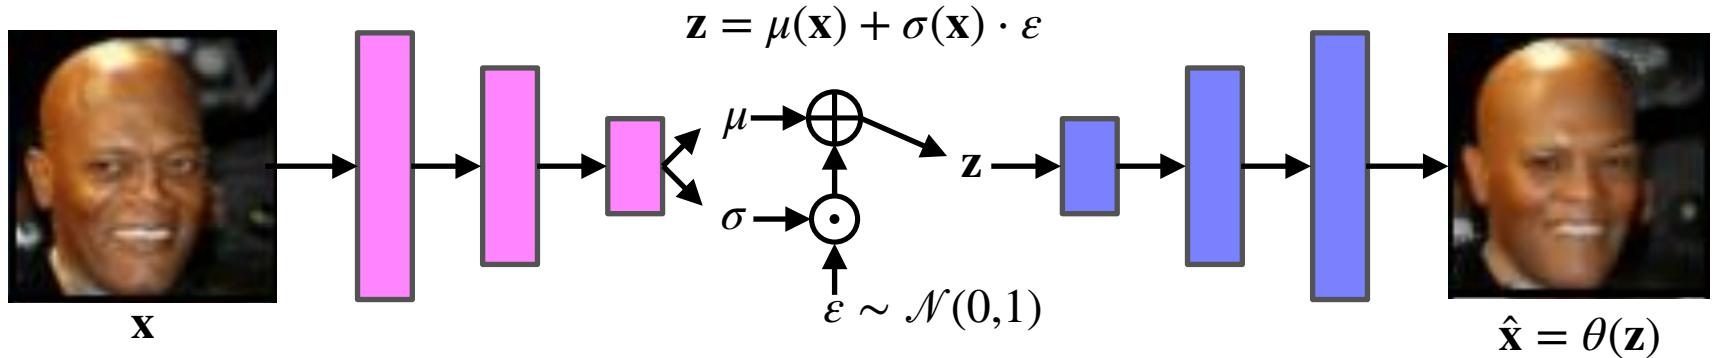
# VANILLA VAE



We approximate expected values using a single sample:

$$ELBO = \underbrace{\ln \mathcal{N}(\mathbf{x} | \theta(\mathbf{z}), 1)}_{\text{RE}} - \underbrace{\left[ \ln \mathcal{N}(\mathbf{z} | \mu(\mathbf{x}), \sigma^2(\mathbf{x})) - \ln \mathcal{N}(\mathbf{z} | 0, 1) \right]}_{\text{KL}}$$

# VANILLA VAE



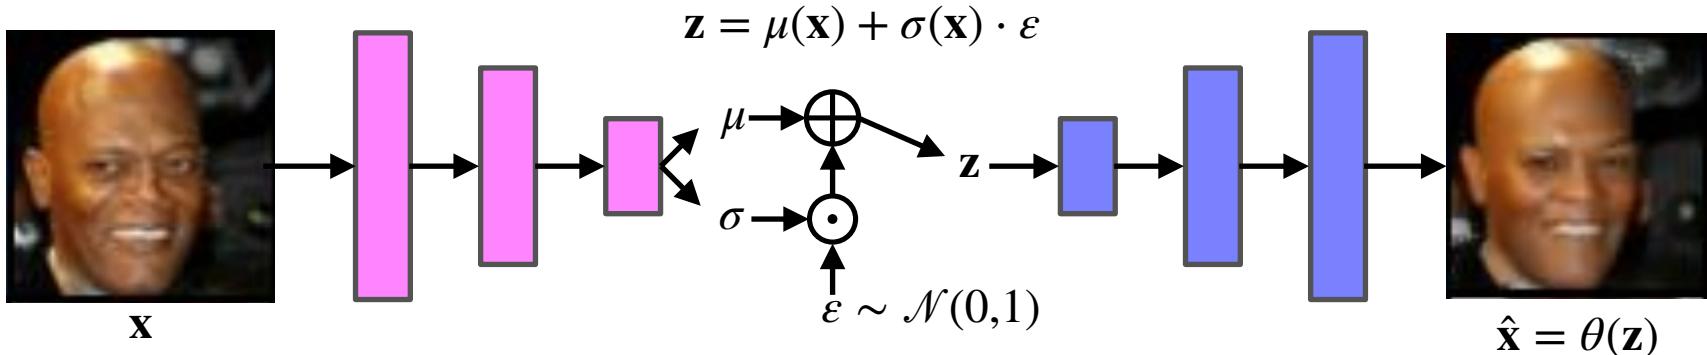
We approximate expected values using a single sample:

**We assume a Gaussian variational posterior.**

$$ELBO = \underbrace{\ln \mathcal{N}(\mathbf{x} | \theta(\mathbf{z}), 1)}_{\text{RE}} - \underbrace{\left[ \ln \mathcal{N}(\mathbf{z} | \mu(\mathbf{x}), \sigma^2(\mathbf{x})) - \ln \mathcal{N}(\mathbf{z} | 0, 1) \right]}_{\text{KL}}$$

**We assume a standard Gaussian prior.**

# VANILLA VAE



We approximate expected values using a single sample:

$$ELBO = \underbrace{\ln \mathcal{N}(\mathbf{x} | \theta(\mathbf{z}), 1)}_{\text{RE}} - \underbrace{\left[ \ln \mathcal{N}(\mathbf{z} | \mu(\mathbf{x}), \sigma^2(\mathbf{x})) - \ln \mathcal{N}(\mathbf{z} | 0, 1) \right]}_{\text{KL}}$$

**REMEMBER! We cannot pick an arbitrary distribution. We must choose a distribution that is appropriate for our data.**

# VANILLA VAE

```
import torch.nn as nn

class VAE(nn.Module):
    def __init__(self, D, M):
        super(VAE, self).__init__()
        self.D = D
        self.M = M

        self.enc1 = nn.Linear(in_features=self.D, out_features=300)
        self.enc2 = nn.Linear(in_features=300, out_features=self.M*2)

        self.dec1 = nn.Linear(in_features=self.M, out_features=300)
        self.dec2 = nn.Linear(in_features=300, out_features=self.D)

    def reparameterize(self, mu, log_std):
        std = torch.exp(log_std)
        eps = torch.randn_like(std)
        z = mu + (eps * std)
        return z
```

# VANILLA VAE

```
def forward(self, x):
    # encoder
    x = nn.functional.relu(self.enc1(x))
    x = self.enc2(x).view(-1, 2, self.M)

    # get mean and log-std
    mu = x[:, 0, :]
    log_std = x[:, 1, :]

    # reparameterization
    z = self.reparameterize(mu, log_std)

    # decoder
    x_hat = nn.functional.relu(self.dec1(z))
    x_hat = self.dec2(x)
    return x_hat, mu, log_std
```

# VANILLA VAE

```
def elbo(self, x, x_hat, z, mu, log_std):
    # reconstruction error
    RE = nn.loss.mse(x, x_hat)

    # kl-regularization
    # We assume here that log_normal is implemented
    KL = log_normal(z, mu, log_std) - log_normal(z, 0, 1)

    # REMEMBER! We maximize ELBO, but optimizers minimize.
    # Therefore, we need to take the negative sign!
    return -(RE - KL)
```

# COMMON ISSUES WITH VAEs

$$q_{\phi}(\mathbf{z}|\mathbf{x}) \propto p_{\theta}(\mathbf{x}|\mathbf{z}) p_{\lambda}(\mathbf{z})$$

Weak **decoders** → bad generations/reconstructions

Weak **encoders** → bad latent representation, *posterior collapse*  
(variational posterior = prior).

Weak **marginals** → bad generations

Variational **posteriors** → what family of distributions?

# COMPONENTS OF VAEs

$$q_{\phi}(\mathbf{z}|\mathbf{x}) \propto p_{\theta}(\mathbf{x}|\mathbf{z}) p_{\lambda}(\mathbf{z})$$



ResNets, DenseNets  
DRAW  
Autoregressive models  
Normalizing flows

# COMPONENTS OF VAEs

$$q_{\phi}(\mathbf{z}|\mathbf{x}) \propto p_{\theta}(\mathbf{x}|\mathbf{z}) p_{\lambda}(\mathbf{z})$$

ResNets, DenseNets

**Normalizing flows**

Hyperspherical dist.

Hyperbolic-normal dist.

Group theory

ResNets, DenseNets

DRAW

Autoregressive models

Normalizing flows

# COMPONENTS OF VAEs

$$q_{\phi}(\mathbf{z}|\mathbf{x}) \propto p_{\theta}(\mathbf{x}|\mathbf{z}) p_{\lambda}(\mathbf{z})$$

ResNets, DenseNets

**Normalizing flows**

Hyperspherical dist.

Hyperbolic-normal dist.

Group theory

ResNets, DenseNets

DRAW

Autoregressive models

Normalizing flows

Autoregressive models

Normalizing flows

VampPrior

Implicit prior

# COMPONENTS OF VAEs

$$q_{\phi}(\mathbf{z}|\mathbf{x}) \propto p_{\theta}(\mathbf{x}|\mathbf{z}) p_{\lambda}(\mathbf{z})$$

ResNets, DenseNets  
**Normalizing flows**  
Hyperspherical dist.  
Hyperbolic-normal dist.  
Group theory

ResNets, DenseNets  
DRAW  
Autoregressive models  
Normalizing flows

Autoregressive models  
Normalizing flows  
VampPrior  
Implicit prior

$$\text{ELBO}(\mathbf{x}; \theta, \phi, \lambda)$$

Adversarial learning  
MMD  
Wasserstein AE

# COMPONENTS OF VAEs

$$q_{\phi}(\mathbf{z}|\mathbf{x}) \propto p_{\theta}(\mathbf{x}|\mathbf{z}) p_{\lambda}(\mathbf{z})$$

ResNets, DenseNets

Normalizing flows

Hyperspherical

Hyperbolic-nor

Group theory

Hierarchical VAEs!

ResN

gressive models

ssive models

Normalizing flows

Normalizing flows

VampPrior

Implicit prior

$$\text{ELBO}(\mathbf{x}; \theta, \phi, \lambda)$$



Adversarial learning

MMD

Wasserstein AE

# HIERARCHICAL VAEs

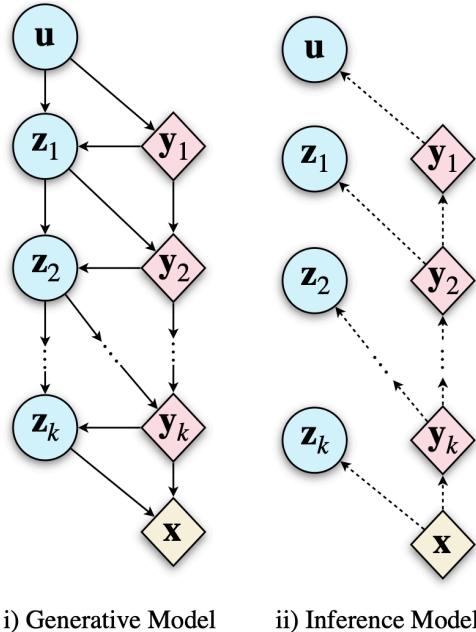


Figure 4: Hierarchical selfVAE.

Gatopoulos, I., & Tomczak, J.M. (2020).  
Self-supervised Variational Auto-Encoders

Vahdat, A., & Kautz, J. (2020).  
Nvae: A deep hierarchical variational autoencoder. NeurIPS 2020

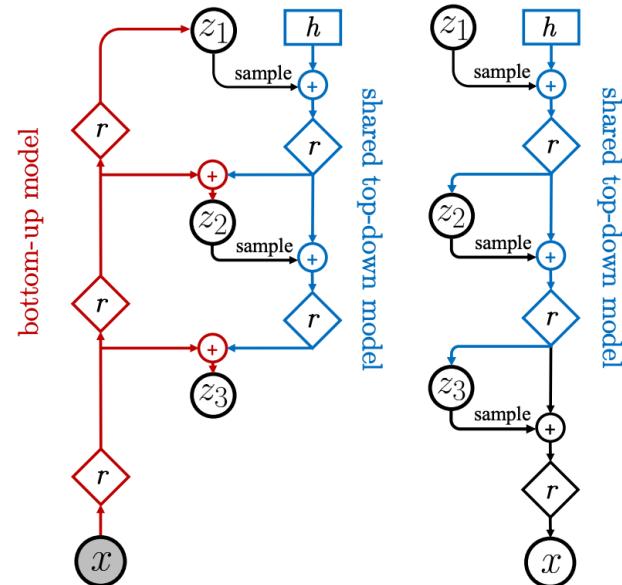


Figure 2: The neural networks implementing an encoder  $q(\mathbf{z}|\mathbf{x})$  and generative model  $p(\mathbf{x}, \mathbf{z})$  for a 3-group hierarchical VAE.  $\diamond^r$  denotes residual neural networks,  $\odot^+$  denotes feature combination (e.g., concatenation), and  $\boxed{h}$  is a trainable parameter.

# HIERARCHICAL VAEs

i) selfVAE - downscale - 3lvi



ii) selfVAE - sketch



iii) VAE - RealNVP



# HIERARCHICAL VAEs

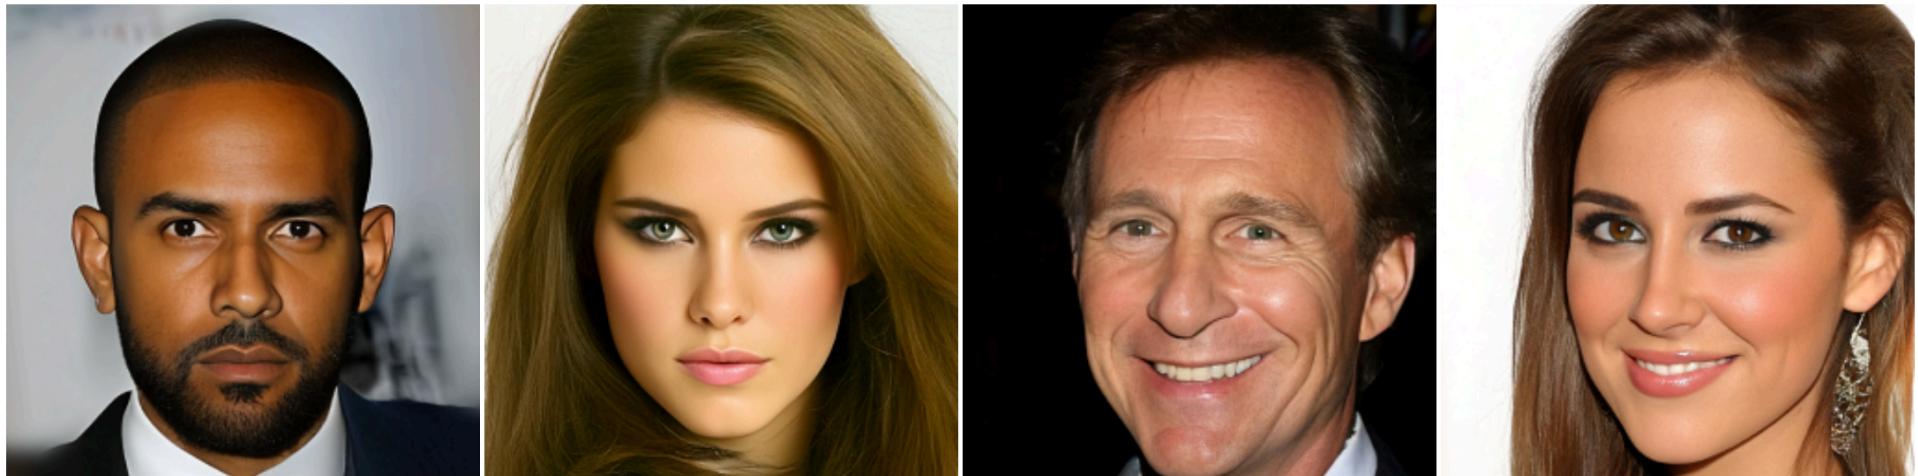


Figure 1:  $256 \times 256$ -pixel samples generated by NVAE, trained on CelebA HQ [28].

## Advantages

- ✓ Non-linear transformations.
- ✓ Stable training.
- ✓ Allows compression.
- ✓ Allows to generation.
- ✓ The likelihood could be approximated.

## Disadvantages

- No analytical solutions.
- No exact likelihood.
- Potential mismatch between true posterior and variational posterior
- Blurry images

## Advantages

- ✓ Non-linear transformations.
- ✓ Stable training.
- ✓ Allows compression.
- ✓ Allows to generation.
- ✓ The likelihood could be approximated.

## Disadvantages

- No analytical solutions.
- No exact likelihood.
- Potential mismatch between true posterior and variational posterior
- ~~Blurry images~~

# Thank you!