

Lecture 11: Deep Reinforcement Learning

Emile van Krieken

Deep Learning 2020

divu.github.io



THIS LECTURE

1. (Deep) Q-Learning
2. Actor-Critic methods
3. Model-based RL (World Models)

2



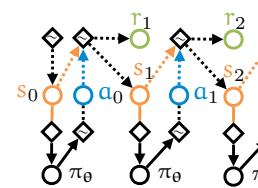
RECAP

REINFORCE:

1. $\tau \sim p(\tau|\theta)$
2. $\theta \leftarrow \theta + \alpha \sum_{t=0}^{T-1} \gamma^t G_t \nabla_\theta \log \pi_\theta(a_t|s_t)$

Discounted reward to-go

$$G_t = \sum_{t'=t}^{T-1} \gamma^{t'-t} r_{t'+1}$$



3



NOTATION UPDATE

- Declutter notation:
- Current timestep t:
 $r_t = r, a_t = a, s_t = s, G_t = G$
- Next timestep t + 1:
 $r_{t+1} = r', a_{t+1} = a', s_{t+1} = s', G_{t+1} = G'$

A problem with RL is that the notation can become very cluttered and hard to understand. To prevent information overflow in the next slides, we will redefine some symbols.

For variables like the **reward**, **action**, **state** and expected **reward** to-go, we just use the symbol without the timestep to refer to the variable at timestep t.

We do the same for the 'next' timestep, t+1, but adding the prime character ('') after the symbol to distinguish it from the current state.

4 VU

Q-FUNCTION

Discounted reward to-go
 $G = \sum_{t'=t}^{T-1} \gamma^{t'-t} r_{t'+1}$

Q-function (state-action value function):

$$Q^\pi(s, a) = \mathbb{E}_{p(\tau|\pi, s, a)}[G]$$

First, we will introduce two very important notions in RL: The **Q-function**. These are defined in terms of the **discounted reward** to-go that we not-coincidentally used in REINFORCE.

We will start by the harder one: The **Q-function**. It is also known as the **state-action** value function, for good reasons! The **Q-function** is the **expected discounted reward** to-go, after executing **action** a in **state** s, **and** afterwards following the **policy** π (!!!!!). I cannot stress this enough, so read it again: After the **action** is executed, the **Q-function** expects the **agent** to follow the **policy** π . It will come as no surprise that we will learn **Q-functions** in deep Q-learning :)

5 VU

OPTIMALITY

Optimal Q function: For all **states** s and **actions** a:

$$Q^*(s, a) = \max_{\pi} Q^\pi(s, a)$$

Given $Q^*(s, a)$, the **optimal policy** is

$$\pi^*(a|s) = \begin{cases} 1 & \text{if } a = \arg \max_{a'} Q^*(s, a') \\ 0 & \text{otherwise} \end{cases}$$

Optimal policy is deterministic!

Next, we will introduce what the **optimal policy** in Reinforcement Learning is. This is the best **policy** that we can perform! It is a **policy** such that, for all **states**, the **value function** is maximized.

What does this **policy** look like? By expanding the **value function**, we see that it is an expectation of the **Q-function** over the optimal **policy**.

Now, we note that the optimal policy is necessarily deterministic. Why does it need to be deterministic? Note that the probability distribution of an expectation needs to sum to 1. Clearly, it is best to put all probability mass on the action that maximizes the Q-function. That is a deterministic policy!

In other words, the optimal policy is the policy that has a probability of 1 for (or all probability mass on) the action maximizing the Q-function.

6 VU

RECURSIVE REWARD TO GO

$$\begin{aligned} G &= r' + \gamma(r_{t+2} + \gamma r_{t+3} + \dots + \gamma^{T-2-t} r_T) \\ &= r' + \gamma \sum_{t'=t+1}^{T-1} \gamma^{t'-t} r_{t'+1} \\ &= r' + \gamma G' \end{aligned}$$

Next, we rewrite the discounted reward to go. Simply write out the sum over discounted rewards. Leave one discount factor out. The sum inside the braces is exactly equal to $G_{\{t+1\}}$!



OPTIMAL Q-FUNCTION

$$Q^*(s, a) = \max_{\pi} \mathbb{E}_{p(\tau|\pi, s, a)}[G] = \mathbb{E}_{p(\tau|\pi^*, s, a)}[r' + \gamma G'] = Q^{\pi^*}(s, a)$$

Expand expectation:

$$\begin{aligned} &= \mathbb{E}_{p(s', r'|s, a)}[r' + \gamma \mathbb{E}_{p(\tau|\pi^*, s')}[G']] \\ &= \mathbb{E}_{p(s', r'|s, a)}[r' + \gamma \mathbb{E}_{\pi^*(a'|s')}[\mathbb{E}_{p(\tau|\pi^*, s', a')}[G']]] \\ \text{That's the Q-function!} \\ &= \mathbb{E}_{p(s', r'|s, a)}[r' + \gamma \mathbb{E}_{\pi^*(a'|s')}[Q^*(s', a')]] \end{aligned}$$

We will now look how to make the definition of the **Q-function** recursive, that is, it can be computed using itself.

We just found a recursive definition for G , so plug that in. The resulting expectation is over a sum over the next **reward** r' and the rest of the **discounted reward** G' . Also, we replace the maximum over policies by the optimal policy. After all, we cannot do better than that! Conveniently, the resulting expression is the definition of the **Q-function** of the optimal policy: The **Q-function** of the optimal policy IS the optimal **Q-function** :

Next, note that the **next reward** is only dependent on the current **state** s and **action** a , and not on the rest of the trajectory. So, we can split the expectation in two: First we sample the next **state** and **reward**, and then the rest of the trajectory.

The first thing that happens in the trajectory then is sampling the next **action** a' . Of course, this happens from the optimal policy π^* . So, expand this expectation.

Now, the most inner expectation is exactly the **Q function** at the next timestep! So let's plug that in.

We now have a recursive definition of the **Q-function**! But we can simplify it even more.

OPTIMAL Q-FUNCTION

$$Q^*(s, a) = \mathbb{E}_{p(s', r'|s, a)}[r' + \gamma \mathbb{E}_{\pi^*(a'|s')}[Q^*(s', a')]]$$

recall:

$$\pi^*(a|s) = \begin{cases} 1 & \text{if } a = \arg \max_a Q^*(s, a') \\ 0 & \text{otherwise} \end{cases}$$

so:

$$Q^*(s, a) = \mathbb{E}_{p(s', r'|s, a)}[r' + \gamma \max_{a'} Q^*(s', a')]$$

No explicit policy in this equation!

Back to that recursive formulation of the optimal **Q-function**. Recall the optimal policy π^* . This is a deterministic policy, that simply chooses the **action** that maximizes the optimal **Q-function**!

This means we can simplify the recursive formulation even further: It chooses the maximizing **action** a' , so we don't need that expectation over **actions**. Simply take the maximum of the optimal **Q-function** here!

This gives the so-called **Bellman optimality equation**: A way to define the optimal **Q-function** without any reference to an external policy.



OPTIMAL Q-FUNCTION

The **optimal Q-function**:

$$Q^*(s, a) = \mathbb{E}_{p(s', r'|s, a)}[r' + \gamma \max_{a'} Q^*(s', a')]$$

Idea: Estimate $Q^*(s, a)$ with NN $Q_\theta(s, a)$!

Policy:

$$\pi_\theta(a|s) = \begin{cases} 1 & \text{if } a = \arg \max_{a'} Q_\theta(s, a') \\ 0 & \text{otherwise} \end{cases}$$

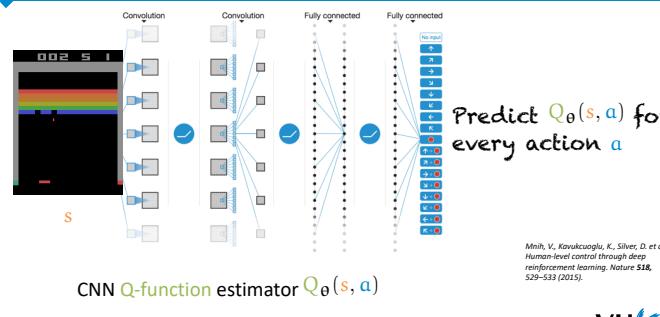
10



We found that the optimal Q-function satisfies this very nice recursion. So what exactly does it say? The optimal Q-function has that it is equal to the expected next reward, plus the discounted optimal Q-function in the next state.

We want to find a Q-function for which this equation holds. Why? From this Q-function we can extract the optimal policy, which takes the action that maximizes the Q-function.

ATARI Q-FUNCTION ESTIMATOR



11



What does a neural network that estimates Q-functions look like? Just like for policies, the neural network usually actually only receives the state, and not the action.

The output layer is important here: It is simply linear! For each action, it outputs a scalar value, which predicts the Q-function. This way, we don't need to run the neural network for each action: We run it once and immediately get Q-function estimations for each action.

BOOTSTRAP ERROR

$$Q^*(s, a) = \mathbb{E}_{p(s', r'|s, a)}[r' + \gamma \max_{a'} Q^*(s', a')]$$

So:

$$Q^*(s, a) - \mathbb{E}_{p(s', r'|s, a)}[r' + \gamma \max_{a'} Q^*(s', a')] = 0$$

Minimize **bootstrapped** error using regression:

$$\arg \min_{\theta} \left(\underbrace{Q_\theta(s, a)}_{\text{prediction}} - \underbrace{\mathbb{E}_{p(s', r'|s, a)}[r' + \gamma \max_{a'} Q_\theta(s', a')]}_{\text{target}} \right)^2$$

12

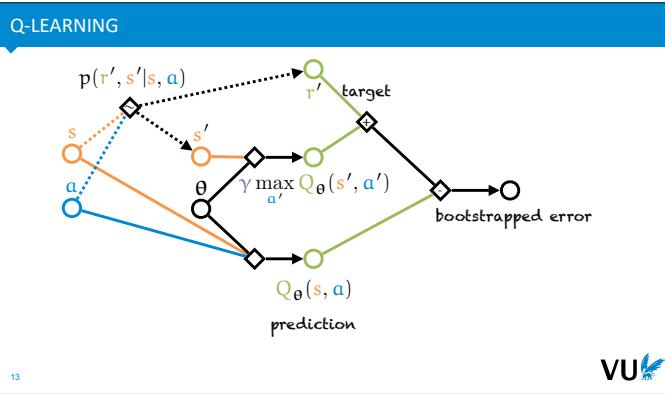


Again, recall the recursive definition of the optimal Q-function. Clearly, subtracting the right-hand side, this definition can be seen as the fact that the difference between the Q-function at the current state minus the second term (expected next reward + discounted Q-function) should be 0.

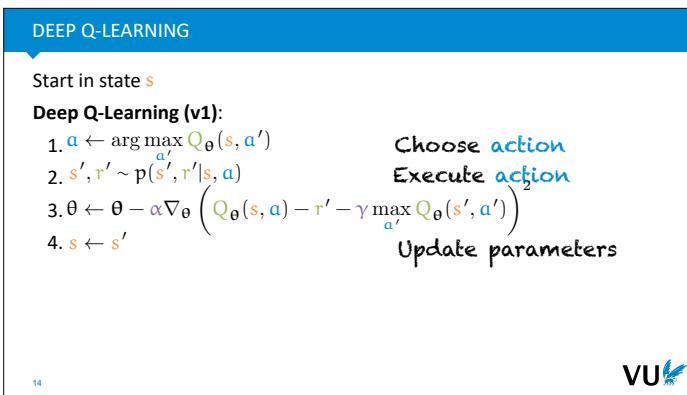
We want our Q-function estimate Q_θ to approximate the optimal Q-function. So, it should also make sure that this equality holds!

The **bootstrapped** error is the difference between our estimate of the Q-function at the current state minus the expected next reward and the maximum of the estimated discounted Q-function. Why is it called bootstrapped? Because the error uses the Q-function estimator itself as a target!

In Deep-Q learning, we minimize the squared bootstrapped error, and use regression techniques to optimize this.



Here we see the computation graph of Q-learning. Compare it closely with the equation on the previous slide to see if you can follow the computation. Note that the same set of parameters is used both to compute the prediction and the target.

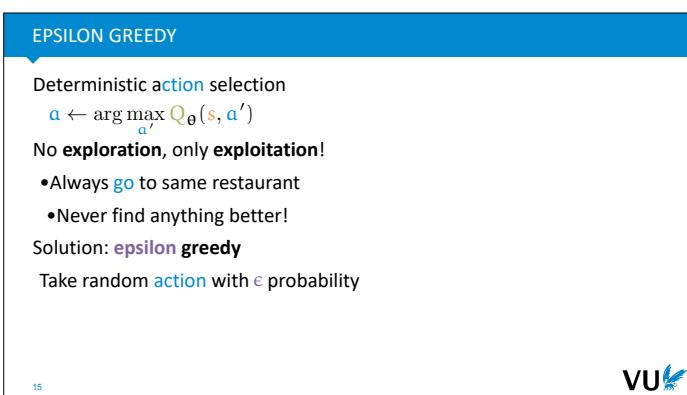


Let's look at our first version of the Deep Q-learning algorithm. This is an **online** algorithm, which means that we update our policy (in this case, Q-value parameters) immediately after executing an **action**, instead of waiting until the trajectory is over.

The algorithm is pretty simple:

1. Choose an **action**. This happens by taking the **action** that maximizes the estimated **Q-value**.
2. Execute the **action** in the **environment**. We transition to the next **state**, and receive a **reward**, just like before.
3. Update the parameters using the bootstrapped error equation. This is a simple gradient descent step.

Repeat this!

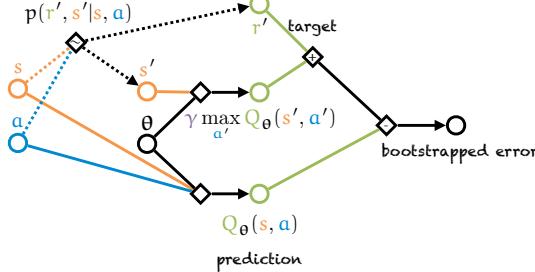


Unfortunately, this is a very, very poor algorithm. One of the reasons is that our **action** selection algorithm is deterministic: We always execute the same **action** in a state (namely, the one maximizing the **Q-value**).

Analogy: Say we want to figure out what the best restaurant in town is. With this way of **action** selection, we'd **take** the first restaurant that seems nice, **exploit** that knowledge, and then always **visit** that restaurant without ever **exploring** any alternatives. Obviously, this leads to sub-optimal behaviour as there's probably a much better **restaurant** in town! This problem is called the **exploration - exploitation tradeoff**.

The most common solution to this problem is **epsilon greedy**. Instead of always taking the best **action**, we simply take a random **action** with **epsilon** probability.

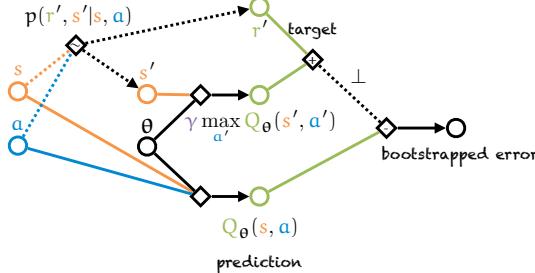
MOVING TARGET



16



SIMPLE SOLUTION



17



DEEP Q-LEARNING

Deep Q-Learning (v2):

1.
$$a = \begin{cases} \text{random action} & \text{with } \epsilon \text{ probability} \\ \arg \max_{a'} Q_\theta(s, a') & \text{otherwise.} \end{cases}$$
 Epsilon-greedy
2. $s', r' \sim p(s', r'|s, a)$
3. $\theta \leftarrow \theta - \alpha \nabla_\theta \left(Q_\theta(s, a) - \perp (r' + \gamma \max_{a'} Q_\theta(s', a')) \right)^2$
4. $s \leftarrow s'$

18



Another issue is the following. Let's see what happens in our backwards pass. How are the parameters θ updated? As you can see, it happens both through the prediction, and *also* through the target **Q-value** prediction! This means that by minimizing the bootstrapped error equation, we change both the prediction and the target. If that seems wrong, from a machine learning perspective, then you're most certainly right.

In fact, if you'd try this out, it'd be very hard to get it to work, as there is no clearly defined target for the regression problem.

A very simple, though somewhat hacky, solution, is to just remove the gradient from the target, again using that weird bot symbol we saw before during the first RL lecture. This already helps quite a bit.

Time to look at the second version of our Deep Q-learning algorithm. Not much has changed, though.

In the first line, we change how we select actions, from the "greedy" always take the best action, to epsilon-greedy, that takes a random action with epsilon probability.

We also added a bot symbol to show we don't put gradients throughout the target of the bellman error.

DEEP Q-LEARNING

Deep Q-learning requires many tricks to get working

- Experience replay
- Target networks
- Double Q-learning
- Multi-step returns
- Gradient clipping
- ...

In assignment 5c, you can explore some of these for bonus points!

19



Unfortunately, Deep Q-learning requires quite a lot of tricks to get working.

The most important one is experience replay. It solves two problems: states are sequentially correlated, and so do not satisfy the identically and independently distributed (iid) assumption that machine learning requires! Experience replay instead saves some states to 'replay' later to train the Q-function. It separates data collection and training the Q-function on it.

There are much more tricks to explore, but we won't get into that here. In assignment 5c, you can play with these!

SUMMARY: DEEP Q-LEARNING

Deep Q-learning

Estimate optimal Q-function

Minimize bootstrapped error

Requires many tricks to get working

Next: Combine ideas in Actor-critic

20



ACTOR-CRITIC



RECAP

Policy gradient:

$$\nabla_{\theta} \mathbb{E}_{p(\tau|\theta)}[R_y] = \mathbb{E}_{p(\tau|\theta)} \left[\sum_{t=0}^{T-1} \gamma^t G_t \nabla_{\theta} \log \pi_{\theta}(a_t|s_t) \right]$$

REINFORCE:

Estimate with Monte Carlo rollout:

$$\nabla_{\theta} J(\theta) \approx \sum_{t=0}^{T-1} \gamma^t G_t \nabla_{\theta} \log \pi_{\theta}(a_t|s_t)$$

High variance :(

22



BASELINES

REINFORCE:

$$\nabla_{\theta} \mathbb{E}_{p(\tau|\theta)}[R_y] = \mathbb{E}_{p(\tau|\theta)} \left[\sum_{t=0}^{T-1} \gamma^t G_t \nabla_{\theta} \log \pi_{\theta}(a_t|s_t) \right]$$

Reduce variance with **baseline** b_t :

$$\nabla_{\theta} \mathbb{E}_{p(\tau|\theta)}[R_y] = \mathbb{E}_{p(\tau|\theta)} \left[\sum_{t=0}^{T-1} \gamma^t (G_t - b_t) \nabla_{\theta} \log \pi_{\theta}(a_t|s_t) \right]$$

baseline

23



A simple important modification is baselines. A baseline is a value that is subtracted from the reward. This is still the unbiased policy gradient! In the next slide, we'll see why.

BASELINES

$$\begin{aligned} & \mathbb{E}_{\pi_{\theta}(a_t|s_t)}[b_t \nabla_{\theta} \log \pi_{\theta}(a_t|s_t)] \\ &= \sum_{a_t} b_t \cancel{\pi_{\theta}(a_t|s_t)} \frac{\nabla_{\theta} \pi_{\theta}(a_t|s_t)}{\cancel{\pi_{\theta}(a_t|s_t)}} \end{aligned}$$

24



So why is subtracting that baseline unbiased? Consider the expectation of the baseline multiplied by the grad-log of the policy.

Write out the expectation and the derivative, just like in the last lecture when we derived the score function.

BASELINES

$$\begin{aligned}
 & \mathbb{E}_{\pi_\theta(a_t|s_t)}[b_t \nabla_\theta \log \pi_\theta(a_t|s_t)] \\
 &= \sum_{a_t} b_t \cancel{\pi_\theta(a_t|s_t)} \frac{\nabla_\theta \pi_\theta(a_t|s_t)}{\cancel{\pi_\theta(a_t|s_t)}} \\
 &= \sum_{a_t} b_t \nabla_\theta \pi_\theta(a_t|s_t) = b_t \nabla_\theta \sum_{a_t} \pi_\theta(a_t|s_t) \\
 &= b_t \nabla_\theta 1 = 0
 \end{aligned}$$

25



Since the baseline doesn't depend on the action, we can move it out of the summation. (and also the gradient, because it's linear).

Now, we have a sum over probabilities with all possible actions. A probability distribution should sum to 1! And the gradient of 1 is 0.

So, the expected value of the baseline multiplied with the grad-log policy is 0, so we can safely subtract it from the Q-function.

WHAT BASELINE?

REINFORCE with **baseline**:

$$\nabla_\theta \mathbb{E}_{p(\tau|\theta)}[R_Y] = \mathbb{E}_{p(\tau|\theta)} \left[\sum_{t=0}^{T-1} \gamma^t (G_t - b_t) \nabla_\theta \log \pi_\theta(a_t|s_t) \right]$$

baseline ?

Value function:

$$V^\pi(s_t) = \mathbb{E}_{p(\tau|\pi, s_t)}[G_t] = \mathbb{E}_{\pi(a_t|s_t)}[Q^\pi(s_t, a_t)]$$

Value function baseline:

$$\nabla_\theta \mathbb{E}_{p(\tau|\theta)}[R_Y] = \mathbb{E}_{p(\tau|\theta)} \left[\sum_{t=0}^{T-1} \gamma^t (G_t - V^\pi(s_t)) \nabla_\theta \log \pi_\theta(a_t|s_t) \right]$$

26



So what value should this baseline be? Usually, we take the **value function**, or **state value function** (compare with the **Q-function**, which is also called the **state-action value function**).

The value function is the expected **reward-to-go** from a **state** when following the policy. It is related to the Q-function as follows: It's the expected **Q-function** when sampling an **action** from the policy.

REINFORCE VS ACTOR-CRITIC

REINFORCE:

I won!



Act, receive reward.

How to reinforce?

$$V^\pi(s_t) = 0.63$$

Increase of random reward wrt expected reward



REINFORCE + baseline:

I won. That result is 37% better than expected!

Random reward

So this reduces the variance, but why?

Consider an agent trying to learn **chess**. Training chess with REINFORCE would do the following: Perform an **action**, then see if in the end we win the game. If so, update based on that: It was a good **action**! Otherwise, if we lost, we say it was a bad **action**. But many things could have happened after taking that **action**, and it could also be very likely that we would lose after taking it.

With that baseline, it's a bit different: again, we reinforce **actions** that lead to winning. But this time, we subtract that by our current belief how likely it is we would win from that **state**. If we did not expect to win, then it turns out that the sequence of **actions** lead to an unlikely **victory**! These **actions** should be reinforced much more than if we were in a **state** that was already looking like it was in a winning position, and then indeed winning.

27

TRAINING VALUE FUNCTION

Like Deep Q-Learning, train neural network V_ϕ with regression.

1. Use rollouts:

$$\Phi \leftarrow \Phi - \alpha \sum_{t=0}^{T-1} (V_\phi(s_t) - G_t)^2 \quad \text{target: reward-to-go}$$

2. Use bootstrapping (lower variance, biased):

$$\Phi \leftarrow \Phi - \alpha \sum_{t=0}^{T-1} (V_\phi(s_t) - \perp(r_{t+1} + \gamma V_\phi(s_{t+1})))^2 \quad \begin{matrix} \text{target: bootstrapped} \\ \text{expected reward-to-go} \end{matrix}$$

28



RECAP

REINFORCE with baseline:

1. $\tau \sim p(\tau|\theta)$
2. $\Phi \leftarrow \Phi - \alpha_c \sum_{t=0}^{T-1} (V_\phi(s_t) - \perp(r_{t+1} + \gamma V_\phi(s_{t+1})))^2$
3. $\theta \leftarrow \theta + \alpha_a \sum_{t=0}^{T-1} \gamma^t (G_t - V_\phi(s_t)) \nabla_\theta \log \pi_\theta(a_t|s_t)$

29



Q-FUNCTIONS IN POLICY GRADIENTS

Recall $Q^\pi(s, a) = \mathbb{E}_{p(\tau|\pi, s, a)}[G]$

Policy gradient:

$$\begin{aligned} \nabla_\theta \mathbb{E}_{p(\tau|\theta)}[R_\gamma] &= \mathbb{E}_{p(\tau|\theta)} \left[\sum_{t=0}^{T-1} \gamma^t G_t \nabla_\theta \log \pi_\theta(a_t|s_t) \right] \\ &= \mathbb{E}_{p(\tau|\theta)} \left[\sum_{t=0}^{T-1} \gamma^t Q^\pi(s_t, a_t) \nabla_\theta \log \pi_\theta(a_t|s_t) \right] \end{aligned}$$

critic actor

Much lower variance!

30

Cool! But we still need to be able to compute that baseline. We'll do this just like in Deep Q-Learning: We'll use a neural network that's going to estimate the **value function**.

There are two options: We can use the **rewards-to-go** from that **state**, just like we'd use for REINFORCE. That's a high-variance target though.

Instead, we usually use bootstrapping errors, just like in Deep Q-Learning: The target is the next **reward** plus the **value function** in the next **state**.

Recall how we defined **Q-functions**: The **expected reward-to-go** G_t from a **state** after performing an **action** (and then continuing to use the corresponding policy).

The **reward-to-go** G_t is the **reward** that is used to update our policy in REINFORCE... Within an expectation too!

Therefore, we can actually replace G_t by the **Q-function**. This gives us an **actor-critic** algorithm, where the **policy** is an actor, and the **Q-function** is the critic: It 'criticizes' taking an **action**.

This has much lower variance! Why? (next slide!)



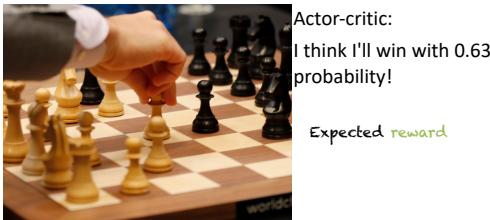
REINFORCE VS ACTOR-CRITIC

Act, receive reward.

How to reinforce?

REINFORCE:

I won!



Random reward

Actor-critic:
I think I'll win with 0.63 probability!

Expected reward

31



BASELINES

Actor-critic:

$$\nabla_{\theta} \mathbb{E}_{p(\tau|\theta)}[R_{\gamma}] = \mathbb{E}_{p(\tau|\theta)} \left[\sum_{t=0}^{T-1} \gamma^t Q^{\pi}(s_t, a_t) \nabla_{\theta} \log \pi_{\theta}(a_t|s_t) \right]$$

Reduce variance even more with value function **baseline**:

$$\nabla_{\theta} \mathbb{E}_{p(\tau|\theta)}[R_{\gamma}] = \mathbb{E}_{p(\tau|\theta)} \left[\sum_{t=0}^{T-1} \gamma^t (Q^{\pi}(s_t, a_t) - V^{\pi}(s_t)) \nabla_{\theta} \log \pi_{\theta}(a_t|s_t) \right]$$

advantage actor

Advantage actor-critic

32



So that's already a very useful modification to the policy gradient (sidenote: Both REINFORCE and actor-critic are called policy-gradient algorithms because (recall last lecture) they estimate the gradient of the expected reward!)

A second important modification is the addition of baselines. A baseline is a value that is subtracted from the critic. This is still the unbiased policy gradient! In the next slide, we'll see why.

ADVANTAGE VS Q-FUNCTION

Act, receive reward.

How to reinforce?

Actor-critic:

I think I'll win with 63% probability!



Expected reward

Advantage actor-critic:
I think I'll be 3% more likely to win.

Expected increase in reward

33



Now compare advantage actor-critic with normal actor-critic. In advantage actor-critic, we don't reinforce based on expected reward, but based on something a bit different:

The difference between expected reward in the state and the expected reward after taking an action! For example, in chess, we reinforce actions of which we think it'll make us 3% more likely to win.

COMPUTING THE ADVANTAGE

Advantage actor-critic:

$$\nabla_{\theta} \mathbb{E}_{p(\tau|\theta)}[R_Y] = \mathbb{E}_{p(\tau|\theta)} \left[\sum_{t=0}^{T-1} \gamma^t (Q^\pi(s_t, a_t) - V_\phi(s_t)) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right]$$

advantage A^π

Estimate V^π with V_ϕ (biased)

Estimate $Q^\pi(s_t, a_t)$ with $r_{t+1} + \gamma V_\phi(s_{t+1})$

$A_\phi(s_t, r_{t+1}, s_{t+1}) = r_{t+1} + \gamma V_\phi(s_{t+1}) - V_\phi(s_t)$

34



So how do we compute the advantage in practice? We can use the **value function** estimation that we introduced when discussing baselines.

We can actually use this estimation to estimate the **Value function baseline** and the **Q-function**, by expanding the **Q-function**. To estimate the **Q-function**, we look what the next **state** is and add the **reward** in the next **state** to the **discounted value function** estimate of that **state**.

It should be noted that this gives biased estimates. This is because we only *estimate* the **value function**, and can definitely make mistakes in its predictions. The policy gradient is only equal to actor critic if **Q** is exactly the **Q-function** of the policy, not for other functions that estimate it! In that sense, actor-critic trades off bias for variance. This is different from REINFORCE with **value function** baselines, which is always unbiased!

In estimating the **Q-function**, there's additional bias by using the **value function** in this way, as it doesn't take the expectation over **actions**.

SUMMARY

Advantage actor-critic:

$$\nabla_{\theta} \mathbb{E}_{p(\tau|\theta)}[R_Y] \approx \sum_{t=0}^{T-1} \gamma^t A_\phi(s_t, r_{t+1}, s_{t+1}) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

where

$$A_\phi(s_t, r_t, s_{t+1}) = r_{t+1} + \gamma V_\phi(s_{t+1}) - V_\phi(s_t)$$

critic baseline

Trade off variance (=sample efficiency) for bias

35



BATCH-MODE ALGORITHM

Advantage actor-critic:

1. $\tau \sim p(\tau|\theta)$
2. $\Phi \leftarrow \Phi - \alpha_c \sum_{t=0}^{T-1} (V_\phi(s_t) - \perp(r_{t+1} + \gamma V_\phi(s_{t+1})))^2$ Update **critic**
3. $\theta \leftarrow \theta + \alpha_a \sum_{t=0}^{T-1} \gamma^t A_\phi(s_t, r_{t+1}, s_{t+1}) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$ Update **actor**

36



We will see versions of advantage actor-critic algorithm, starting with the batch-mode algorithm.

Here, in a loop, we sample a trajectory from the MDP, then update the parameters.

First, we update the critic. We just have to train the **value function** here, as the Q-function is expressed in terms of the value function! Training it happens just like in REINFORCE with a baseline using bootstrapping.

Then we update the actor based on the advantage function.

ONLINE ACTOR-CRITIC

Online Actor-Critic:

1. $a \sim \pi_\theta(a|s)$ Select actions according to policy
2. $s', r' \sim p(s', r'|s, a)$
3. $\Phi \leftarrow \Phi - \alpha_c(V_\Phi(s) - \perp(r' + \gamma V_\Phi(s')))^2$ Update critic
4. $\theta \leftarrow \theta + \alpha_a A_\Phi(s, r', s') \nabla_\theta \log \pi_\theta(a|s)$ Update actor
5. $s \leftarrow s'$

37



This algorithm has a second version, an ‘online-mode’ version. This happens just like in the algorithm of Deep Q-learning. Namely, that don’t sample a complete trajectory, but update each time after taking an action.

It is not much harder: In a loop, we sample an action from the policy and execute it in the environment. Note that we don’t need to do something like epsilon-greedy here because actor-critic already uses stochastic policies!

Now, we simply use the resulting next state and reward to compute an update for a single transition. Again, first the critic, then the actor. Then we assign the next state to the current state and continue looping.

It’s a detail, but the actor update in this algorithm introduces bias compared to the previous algorithm. Can you spot why? (This is not important for the exam or anything, just interesting :))

A2C

Uses only a single sample

And batching over time would give correlated minibatches

A2C: Multiple online agents synchronously act

Collect experiences at each step for minibatch.

Efficient method!

38



VARIANCE REDUCTION

Techniques to reduce variance:

- Baselines:
 - Reinforce difference with expected value
 - Unbiased, fairly low variance
- Actor-critic:
 - Reinforce expected value
 - Biased, fairly low variance
- Advantage actor-critic:
 - Reinforce increase in expected value
 - Biased, low variance

39



ACTOR-CRITIC

Actor-critic methods

- Model **actor**: Policy NN
- Model **rewards**: Value function NN

What about 3rd RL component: **environment**?

Actor-critic methods model two out of three RL components. The actor models **actions** through policy Networks, while **rewards** are modelled through the critic, in our case by estimating the **value function**.

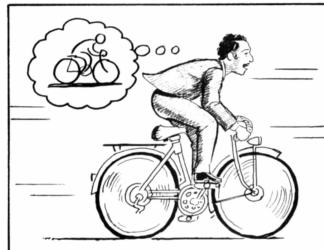
So what about the 3rd RL component: **states** from the **environment**? Can those also be modelled?

41

WORLD MODEL

World Models

- Model **environment** using neural networks!
- Find good state representations



McCloud, Scott. *Understanding Comics: The Invisible Art*. Tundra Publishing, 1993.
Ho, David, and Jürgen Schmidhuber.
"World models."

42

The answer is yes! Modelling the environment using neural networks is often called **world models**. The goal of world models is to find representations that describe the current state well. This can be in multiple ways: It can compress and abstract information, or it can be used to represent knowledge about the world that isn't directly in the current state.

For example, if you have a robot that can look around, it only sees a part of the world. You'd also want to model what is outside of what it can see!

WORLD MODEL

World Models

- Model **environment!**

Model-based RL using Neural Networks

Use **generative model** to learn

1. how to represent **states**

2. to what **state** we transition after taking **action**

3. what **reward** we receive in **state**

43



World Models fall within the field of model-based reinforcement learning, by using neural networks to model the environment.

We use generative modelling (techniques from the last weeks lectures) to create world models.

We have multiple models:

1. A model that learns how to represent individual **states**
2. A model that learns how to transition from one **state** to another after taking an **action** (=transition probability)
3. A model that models the **reward** we receive in a **state**

WORLD MODEL

World Models combine many ideas from this course:

- Deep generative modelling
- Variational Auto-Encoders
- Gradient estimation
- Actor-critic

44



World models are also interesting because they combine a lot of ideas of this course and show well how they fit in the deep learning toolkit!

WORLD MODEL COMPONENTS

Components of our World Model:

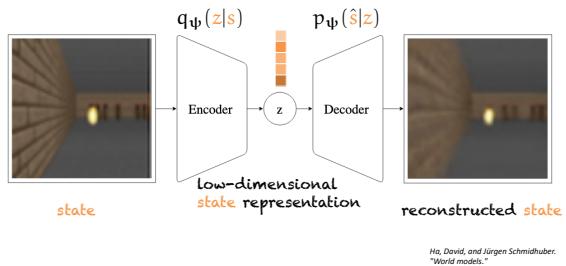
- Reward model $P_\Psi(r|z)$
- Transition model $P_\Psi(z'|z, a)$
- State encoder $q_\Psi(z|s)$
- State decoder $p_\Psi(s|z)$

Hafner, Daniil, et al. "Dream to Control: Learning Behaviors by Latent Imagination." International Conference on Learning Representations. 2019.



Here, we summarize all components of the world model: A model for **rewards**, a model for **state** transitions, and a model for **state representations** (both encoding and decoding the **state representations**). We will start with that last one first. We use a VAE for this.

VARIATIONAL AUTO-ENCODER

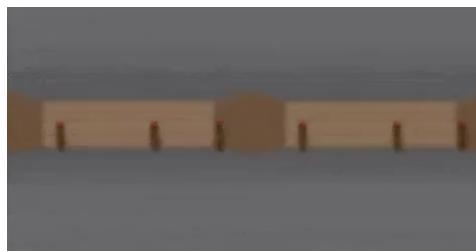


46

We will use a variational autoencoder (VAE) to represent states. For this, we use the latent space of VAEs. This latent space is useful, because the **state representations** allow reconstructing the original **state**. Therefore, they are a low-dimensional representation of the original **state**!

In this case, we'll be looking at **states** represented as images.

VARIATIONAL AUTO-ENCODER



47

Since we have a VAE, we can encode a state (=image), then decode it again to get a reconstruction. This gives a visual idea about how well we have learned to represent the states.

This simple example shows this can work pretty well!

TRAIN VISION MODEL

Training VAE to represent **states**:

1. Collect experience by **acting** in **environment**
2. Train VAE to reconstruct **states**

Goal: Learn useful **representation** z_t of **states** s_t

Compress s_t to what's essential for **acting**



So how do we train a VAE to represent **states** in the **environment**? It's not so hard: We collect our 'training data' by just **acting** in the **real environment**! Then we use this training data to teach the VAE how to reconstruct **states** using the ELBO loss as explained in the VAE lectures.

The goal is to learn useful representations of **states**. But how is useful defined? In the end, all we need to do is to be able to use the **representation** to **act** well in the **environment**.

An interesting question is whether the VAE **representation** is useful. It has to save a lot of information: Namely, everything needed to properly reconstruct an image! But our **agent** ideally gets just a **representation** that retrieves from the **state** was is essential for the agent to **act**. A common discussion point against using VAEs here is that the **representation** needs to do too much: It should both reconstruct **images**, and be easy to **act** with!

48

WORLD MODEL COMPONENTS

Components of our World Model:

- Reward model $P_\Psi(r|z)$
- Transition model $P_\Psi(z'|z, a)$
- State encoder $q_\Psi(z|s)$
- State decoder $P_\Psi(s|z)$

Hafner, Danijar, et al. "Dream to Control: Learning Behaviors by Latent Imagination." International Conference on Learning Representations, 2019.

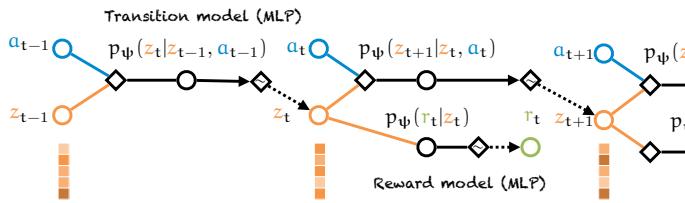


49

We have seen how to encode and decode states to get a **state representation**. Now let's look at the other two models and how they are trained. It should be noted that the **reward** model and transition model, together with **action** inputs, create a markov decision process! We will see why next.

The reward model is a univariate Gaussian distribution, while the transition model is a multivariate Gaussian distribution.

TRANSITION DYNAMICS



Let's look at the detailed computation graph formed using the world model. We start with some **state representation** in timestep $t-1$, in which we take an **action**. We use the transition model to determine the distribution over possible next **states**. From this distribution, we sample the next **state**! Note that this is the **state representation**, not the original state (for example an image). Note that the transition model is a neural network, in this case an MLP!

From this newly generated **state representation**, we use the **reward** model, also an MLP, to predict the **reward** we get in this **state**.

We also use the **state representation**, together with an **action**, to determine the next **state** transition probability. Note that here, we make no assumptions on what **actions** to take, or how we choose these **actions**!

This loop keeps repeating, just like in MDPs. In general, be sure to compare this to MDPs: World Models use almost the same probability distributions as in MDPs except that they model them using neural networks instead of using the **environment** itself. The only difference is that the **environment** generates the next **state** and **reward** together, instead of it being a separate model like in world models.

WORLD MODEL COMPONENTS

Components of our World Model:

- Reward model $P_\Psi(r|z)$
- Transition model $P_\Psi(z'|z, a)$
- State encoder $q_\Psi(z|s)$
- State decoder $P_\Psi(s|z)$

Hafner, Danijar, et al. "Dream to Control: Learning Behaviors by Latent Imagination." International Conference on Learning Representations, 2019.



51

TRAIN DYNAMICS MODEL

1. Collect experience by **acting** in **environment**

2. **foreach** (s, a, r', s')

 1. $z \sim q_\Psi(z|s)$ VAE Encoder

 2. $z' \sim q_\Psi(z'|s')$

 3. Maximize $\log p_\Psi(z'|z, a)$ Transition model

 4. Maximize $\log p_\Psi(r'|z')$ Reward model



Hafner, Danijar, et al. "Dream to Control: Learning Behaviors by Latent Imagination." International Conference on Learning Representations. 2019.



52

Next goal: How do we train the dynamics of the world model? This too can be done by collecting experiences in the real **environment**. Simply have an **agent act** in the real **environment** to collect the training data.

Experiences can be seen as a list of transitions from a **state**, where we take an **action** to get to a new **state** with a **reward**.

To train the world model, we first have to encode the **states** to the latent **state representations** z . Training the transition model is then easy: simply make sure that the transition model predicts the new **state representation** z' from the current **state representation** z . Similarly, the **reward** model is trained to predict the real **reward** from the new **state representation**.

DREAMING

World models can be used to **dream** trajectories

(more formally, **latent imagination**)

Use to train RL agents $\pi_\theta(a|z)$ **without interaction** with **environment**



Hafner, Danijar, et al. "Dream to Control: Learning Behaviors by Latent Imagination." International Conference on Learning Representations. 2019.



53

Now let's get to the exciting part: Dreaming! Because World Models are deep generative models, we cannot just train them, we can also generate new data! Within World Models, this is informally called 'dreaming', and more formally 'latent imagination'. Dreaming basically means imagining trajectories of **states**.

Why is this so useful? It allows us to train **policies** **without** any **interaction** with the **environment** at all. The world models create the **environment** to train on, which means we need way fewer interactions with the environment.

We would often like to minimize how often we **interact** with the **environment**: This might be because the **environment** takes a long time to run, or because it is expensive in some way to **interact** with the **environment** (for example, consider robots!).

The **agents** we train on the world model are trained on the **state representations** z instead of the **states** s themselves.

DREAMING

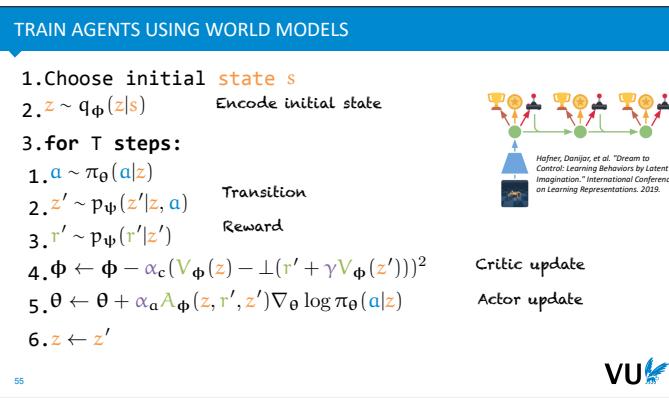


Ho, David, and Jürgen Schmidhuber.
"World models."

We can dream up complete segments of a video game. Here, we start off from some real state, then keep using this latent imagination to dream up a sequence. We can even use our own input here! To play the game 'in the dream', visit <https://worldmodels.github.io/>

54





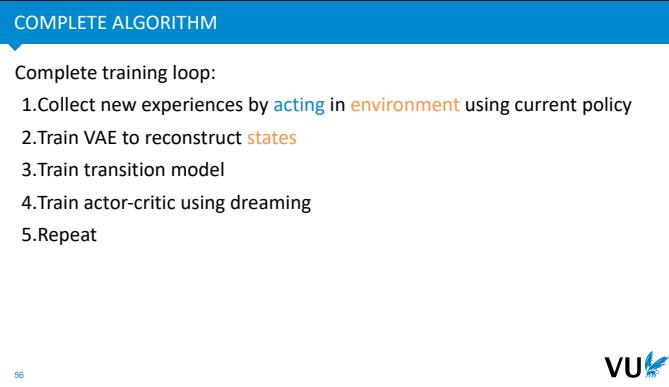
We will adopt an actor-critic online-mode algorithm to train on world models. Assume we have a trained world model, how do we train the **agent**?

We start at an initial **state**, which we encode to the latent **state representation** z . Then we loop: use the policy to choose an **action** based on z .

Use the **action** and z to transition to the next **state representation** z' . Also predict the **reward** in that **state**.

Then we do an update. These are just like in normal (online-mode) advantage actor-critic, except that it uses the **state representations** instead of the **state** itself.

Then we loop: Set the next **state representation** to the current **state representation** z .



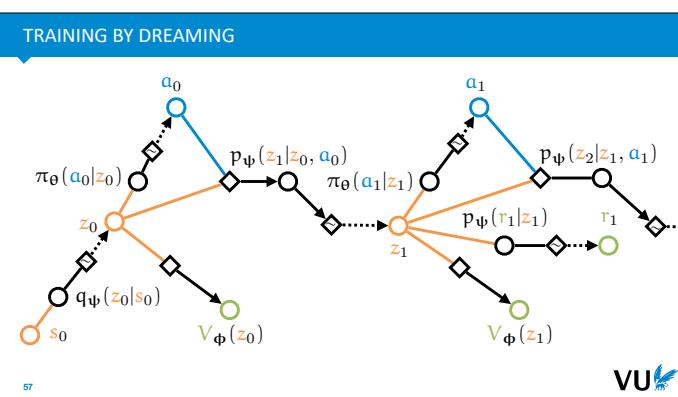
We can now train the VAE and the transition model from experiences in the real **environment**, and train an **agent** from the resulting world models.

Now how does the whole training algorithm look like? From some initialization, we use the current policy to **act** in the **environment**. Note that, as the policy uses **state representations** instead of the **state** itself, so we first have to encode **states** using the VAE encoder.

Next, we train the VAE to reconstruct **states** and the transition model to train the transition dynamics using the new experiences.

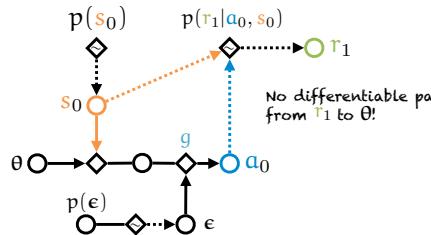
Finally, with the resulting world model, we train the actor-critic **agent**.

Importantly, we repeat this multiple times! Why? Initially, our **agent** is probably pretty bad, and will not be able to explore **states** that are hard to reach. After the **agent** has trained a bit, it's able to reach those new **states**. But now the world model has no knowledge of such **states** yet, and has to learn how to represent these new **states**!



Here, we walk through the computation path when training an actor-critic agent inside a dream.

RECALL: REPARAMETERIZATION IN RL

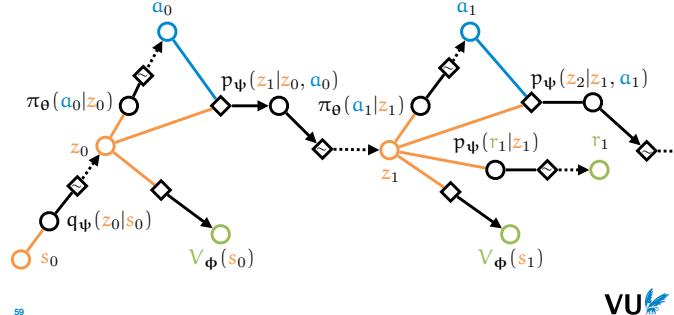


58



Recall the discussion on reparameterization from the last part of the previous lecture. We mentioned that we cannot use the low-variance gradient estimation trick reparameterization in RL because the **environment** is not differentiable. So even if we can make sampling differentiable through reparameterization, this doesn't solve the non-differentiability of the **environment**.

TRAINING BY DREAMING

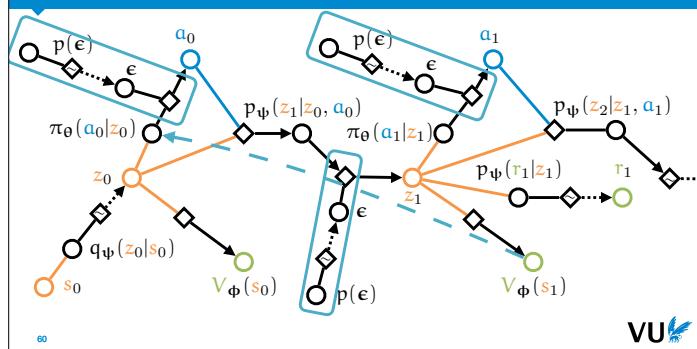


59



Now let's consider the computation graph when dreaming instead of running the true MDP. Notice how relatively few dotted lines there are? This is because we made the **environment** differentiable by modelling it with neural networks! The only non-differentiable steps are sampling, and we know just the trick to make this differentiable :)

TRAINING BY BACKPROPAGATING THROUGH DREAM



60



Here, we modified the computation graph to use reparameterization at all sampling steps. Importantly, we can now follow the graph backward from the **value** functions to the parameters θ of the policy, because everything is differentiable! This means we cannot just dream our **state**, but also 'backpropagate' through the dream to efficiently optimize in it.

TRAIN AGENTS BY BACKPROPAGATING THROUGH WORLD MODELS

```

1. Choose initial state  $s$ 
2.  $z \sim q_\Phi(z|s)$ 
3. for T steps:
   1.  $a \sim \pi_\theta(a|z)$ 
   2.  $z' \sim p_\Psi(z'|z, a)$ 
   3.  $r' \sim p_\Psi(r'|z')$ 
   4.  $\Phi \leftarrow \Phi - \alpha_c(V_\Phi(z) - \perp(r' + \gamma V_\Phi(z')))^2$ 
   5.  $\theta \leftarrow \theta + \alpha_a \nabla_\theta V_\Phi(z')$ 
   6.  $z \leftarrow z'$ 

```

Hafner, Danijar, et al. "Dream to Control: Learning Behaviors by Latent Imagination." International Conference on Learning Representations, 2019.

VU

By just ensuring we reparameterize the samples of the **action** and the **state** transitions, we can do backpropagation through the **value** function to train the **agent**! This means that the **agent** gets very accurate and informative gradient information, instead of the much less informative policy-gradient based methods.

DREAMER

Hafner, Danijar, et al. "Dream to Control: Learning Behaviors by Latent Imagination." International Conference on Learning Representations, 2019.

VU

WORLD MODELS

- Keep a model of **environment**
- VAE represents **states**
- Transition model represents dynamics
- Train **agent** using world models
- “Dream” trajectories
- No **interaction** with **environment**!
- Backpropagate through dreams

VU